

Guide for Advanced Algorithms for Australia and New Zealand Algorithmics & Computing League Competition.

August 27, 2012



This work is licensed under the Creative Commons Attribution-ShareAlike 3.0 Australia License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-sa/3.0/au/>.

Prepared By: Darran Kartaschew

Document Version: 1.0

Last Updated on: August 27, 2012

Contents

1	Introduction.....	1
1.1	About the Competition	1
1.1.1	ANZAC 2012.....	1
1.1.2	ACM-ICPC.....	1
1.2	About this Guide	1
1.2.1	Development System	2
1.3	License	3
2	Supported Competition Environments.....	4
2.1	PC ^ 2.....	4
2.1.1	Submit Run.....	5
2.1.2	Submit Test.....	6
2.2	Software Languages	6
2.2.1	Java	6
2.2.2	C++	7
2.2.3	C#	7
2.3	IDEs.....	7
2.3.1	Eclipse	7
2.3.2	Visual Studio.....	9
2.4	Submission Guidelines	13
3	Performance	14
3.1	BigO Notation	14
3.2	Measuring Performance.....	14
3.2.1	C#	15
3.2.2	C++	15
3.2.3	Java	15
3.2.4	Integers vs Floating Point.....	15
3.3	Implementation and Modern Software Engineering Practices.....	17
4	Basic Source Templates.....	18
4.1	Input / Output.....	18
4.1.1	C#	18
4.1.2	C++	20
4.1.3	Java	22
5	Basic Algorithms.....	24
5.1	Sorting	24
5.2	Searching	24
5.2.1	Linear Search.....	24
5.2.2	Binary Search.....	26
5.3	Array Handling.....	29
5.3.1	Array Performance	29
5.3.2	Array Traversal Methods.....	32
5.3.3	Diagonal Traversal of an Array	33
5.3.4	Array Rotation	35
5.3.5	Array Mirroring or Flipping	38
6	Advanced Algorithms.....	40
6.1	Simple Maths	40
6.1.1	Greatest common divisor	40
6.1.2	Sieve of Eratosthenes (prime number generation).....	42
6.2	String based algorithms and data structures	48
6.2.1	Brute Force Substring Search.....	48
6.2.2	Knuth-Morris-Pratt Substring Search	55
6.3	Graph Theory (Basic)	60
6.3.1	Representation as adjacency matrix and adjacency list	61

A	References	66
---	------------------	----

List of Figures

2.1	PC ^ 2 Login Screen	4
2.2	PC ^ 2 Client Submit Run Tab	5
2.3	Sample Judge's Response	6
2.4	Eclipse - New Java Class	8
2.5	Eclipse - New Java Application	9
2.6	Visual Studio C# New Project	10
2.7	Visual Studio C# Program.cs	10
2.8	Visual Studio C++ New Project	11
2.9	Visual Studio C++ Project Options	12
2.10	Visual Studio C++ Program.cpp	12
5.1	Diagonal Traversal of an Array with a Top-Left Origin	34
5.2	Array Rotation	36
6.1	Demonstration of Sieve of Eratosthenes	43
6.2	Example of Brute Force String Matching	49
6.3	Scanning Grid for Where's Waldorf	54
6.4	Illustration of the KMP algorithm.	56
6.5	Failure Function for Figure 6.4	56
6.6	Basic Graph	60
6.7	Basic Graph Example (weighted, undirected)	61

List of Algorithms

1	Linear Search	24
2	Binary Search	27
3	Row-wise Traversal of an Array	32
4	Column-wise Traversal of an Array	32
5	Diagonal Traversal of an Array	33
6	Array Rotation Clockwise	35
7	Array Rotation AntiClockwise	35
8	Mirror Array Along Vertical Axis	39
9	Mirror Array Along Horizontal Axis	39
10	Euclidean Algorithm (Iterative)	40
11	Euclidean Algorithm (Recursive)	40
12	Sieve of Eratosthenes	43
13	Brute Force String Matching	49
14	Knuth-Morris-Pratt Substring Search	55

List of Source Code Implementations and Solutions

1	Timing - C#	15
2	Timing - C++	15
3	Timing - Java	15
4	Source code for Timing test	16
5	C# Input Example	19
6	C++ Input and Output Example	20
7	C++ getline() example	20
8	C++, cout vs printf()	22
9	Java Input and Output Example	23
10	Linear Search Implementation (Java)	25
11	Linear Search Implementation (C++)	25
12	Solution to Linear Search (Java)	26
13	Binary Search (Java)	27
14	Solution to Binary Search Problem (Java)	28
15	Array Access Performance (C++)	31
16	Diagonal Traversal of an Array (Java)	34
17	Array Clockwise Rotation (Java)	36
18	Array Anti-Clockwise Rotation (Java)	36
19	Solution to Array Rotation Problem (Java)	37
20	Euclidean Algorithm (Java)	41
21	Solution to GCD Problem (Java)	41
22	Sieve of Eratosthenes Algorithm (Java)	44
23	Solution to Primes for Hashtable (Java)	46
24	Brute Force String Matching (C++)	50
25	Solution to Where's Waldorf (C++)	51
26	KMP String Match (C++)	57
27	Solution to Big String Search (C++)	58
28	Adjacency List (Java)	62
29	Solution to Adjacency List Problem (Java)	63

1 Introduction

1.1 About the Competition

The programming contests held in Universities across Australia and New Zealand, are part of the Australia and New Zealand Algorithmics & Computing League Competition and is used in conjunction with the ACM-ICPC competition. These competitions are aimed at challenging students in completing a set number of problems within the allocated time slot (typically 5 hours), with the winners in each location given some prestige.

In recent years teams from not only Universities taken part, but teams from TAFE and other educational institutions have taken part in the competition. Additionally teams outside of Australia and New Zealand such as those from the Phillipines have also taken part.

1.1.1 ANZAC 2012

The ANZAC 2012 competition takes place in 5 to 6 rounds each year and are sponsored by a local University and associated Faculty member. Typically, a single round will run for 5 hours (starting at midday for East Coast Australia), and at least 6 problems will be presented for completion by students.

All challenges require some form of problem solving skills or techniques and do require at least a basic understanding of different algorithms in order to complete the challenges, let alone to be competitive in the competition.

In order to compete within the competition it is recommended that 3 students form a team to work together on solving the challenges. Each team is only given 1 computer to work on, and all reference material brought into the competition must be in printed form only¹.

Scores are awarded for completed challenges (typically 1 point), and the time elapsed from the start of the competition to accepted submission of the challenge is also noted. If a submitted challenge fails, then a 20 minute time penalty is added to the teams total time value.

As a minimum each contest will allow either C/C++ and Java, however additional programming languages may also be included. Typically C# has been allowed in recent years, due to the popularity of the language, especially as it is taught fairly early in a students undergraduate degree.

Overall, the competition is designed to be challenging, fun and also students to advance within their field of study. It is also a great way to network amongst other equally capable students within the programming field.

1.1.2 ACM-ICPC

The ACM-ICPC competition is an International level competition sponsored by IBM, ACM and Upsilon Pi Epsilon, and contestants who make the world finals are often sort after by industry for later employment, as well as bringing notoriety and prestige to the University or College to which the contestants originate from. The regional component of the competition is typically held as the last ANZAC competition, as both competitions share the same tools, resources and rules.

The top two teams from each region (and in the case of Australia and New Zealand, the top team from Australia and top team from New Zealand), attend the International competition held annually in late March/early April in an overseas location. The 2012 ACM-ICPC Finals consisting of teams from all over the world was held in Warsaw, Poland.

1.2 About this Guide

This guide is designed to give students some background knowledge of the environments utilised within the competition, as well as information on various algorithms needed to solve problems. The included

¹The printed material requirement is to ensure that no copying of existing source code is allowed, only transcription of source code from written form

algorithms are by no means exhaustive, however represent the bulk of the algorithms that will be useful in completion of challenges.

This guide book is split into multiple parts:

1. Basic Source Templates that cover the basic frameworks needed for challenge submissions.
2. Basic Algorithms and techniques.
3. Advanced Algorithms.

All algorithms described will include:

1. A short statement on the algorithm and the intended uses, as well as other possible uses.
2. The pseudocode for the algorithm.
3. An actual implementation in at least 1 programming language. This will typically be in the form of a function or method call.
4. An example challenge that requires the use of the algorithm.
5. An example solution to the challenge.

Throughout the guide there will be notes on performance aspects of each algorithm, as well as helpful utility functions to make better use of the algorithm implementations. One example will be a function to convert an Adjancy List into an Adjancy Matrix used for different graph based algorithms.

1.2.1 Development System

The applications and source code snippets developed for this guide were performed on the following hardware and software combinations as noted below. Any performance measurements, in particular times required for certain functions reflect times as acquired with listed hardware and software combinations. Performance measurements will vary accordingly with different hardware and software combinations when performing your own performance measurements.

1.2.1.1 Hardware

HP xw4600 Workstation, with:

- Intel Core 2 Quad, Q9400 @ 2.66GHz (Quad core, 2.66GHz, 6MB L2 cache, 64bit enabled).
- 4GB RAM (4 × 1GB Reg ECC DDR2-800Mhz)
- 250GB 7200rpm HDD
- nVidia Quadro FX580 graphics card.
- Dual 20" LCD Monitors (1680x1050 resolution).

1.2.1.2 Software

Oracle Solaris 11 11/11, with:

- Solaris Studio 12.3 (C++)
- Netbeans 7.2 (Java)
- Java 6 JDK 1.6u26 or Java 7 JDK 1.7u5
- gcc 4.5.2 (C++)

Microsoft Windows XP x64, with:

- Microsoft Visual Studio 2010 (C# and C++)
- .NET Framework 2.0

1.3 License



This work is licensed under the Creative Commons Attribution-ShareAlike 3.0 Australia License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-sa/3.0/au/>.

This means you are free:

- to copy, distribute, display, and perform the work
- to make derivative works
- to make commercial use of the work

Under the following conditions:

- **Attribution** — You must give the original author credit.
- **Share Alike** — If you alter, transform, or build upon this work, you may distribute the resulting work only under a licence identical to this one.

With the understanding that:

- **Waiver** — Any of the above conditions can be waived if you get permission from the copyright holder.
- **Public Domain** — Where the work or any of its elements is in the public domain under applicable law, that status is in no way affected by the license.
- **Other Rights** — In no way are any of the following rights affected by the license:
 - Your fair dealing or fair use rights, or other applicable copyright exceptions and limitations;
 - The author’s moral rights;
 - Rights other persons may have either in the work itself or in how the work is used, such as publicity or privacy rights.
- **Notice** — For any reuse or distribution, you must make clear to others the licence terms of this work.

Original Author endorsed waivers:

- The original author however allows use of source code snippets, that is, source code written in the languages of C++, C# or Java contained with this guide for any purpose, without attribution. This waiver does not extend to the text, nor other materials contained within the guide.

2 Supported Competition Environments

The guide will focus on Java being developed in Eclipse, and C# being developed in Visual Studio. However there will be examples in C++ when appropriate. Most other IDEs have similar options, when used for development, debugging and/or profiling.

2.1 PC²

The primary tool that allow students to submit their challenge entries to be judged in the PC² Software Suite. The application itself is developed by California State University, Sacramento for the purposes of programming competitions and has been adopted by both the Australia and New Zealand Algorithmics & Computing League (ANZACL) and ACM for their respective competitions.

An example of the Login Interface is shown in Figure 2.1.



Figure 2.1: PC² Login Screen

Once logged into the system, the following options are typically available:

Submit Run Allows you to submit a challenge entry to be judged, or alternatively to test your entry against some supplied sample data.

View Runs Allows you to view a history of submissions made to the judges.

Request Clarification Allows you to request a clarification from the judges about one of the challenges.

View Clarifications Allows you to see the responses to your requests for clarifications.

Options Allows you to access various options that control the clients operation. However this tab, only allows you to view the operational log of the client.

Most of the operations on the various areas are self explanatory, so won't be covered in detail. The main screen that competitors will utilise is the **Submit Run** tab as shown in Figure 2.2.

This screen has two main modes of operation, allow a competitor to test their submission against some sample input, or submit their source code to be judged. Both have similar operations, except the test has one additional step.



Figure 2.2: PC ^ 2 Client Submit Run Tab

2.1.1 Submit Run

To submit a run for judging, perform the following steps:

1. From the Problem dropdown list select the challenge that you are attempting.
2. From the Language dropdown list select the programming language in the submission is written in.
3. Use the Select button to select the source code file for the submission. (Note: A single Source Code file is required, do not attempt to submit data files or executable files).
4. Use the Add Button to select any additional files needed to complete your submission. (Note: This is rarely needed).
5. Click on Submit, and Yes to confirm to have your submission judged.
6. You will receive a confirmation dialog confirming that your entry has been submitted.

Once your entry has been judged you will receive one of the following confirmations:

- Yes - Your submission was successful in passing all tests. Congratulations, you have been awarded one point.
- No - Your submission failed one or more tests.
- Time Overrun - Your submission took more time than allowed for the challenge.

An example Judge's Response Dialog is shown in Figure 2.3.




Figure 2.3: Sample Judge's Response

2.1.2 Submit Test

Before you submit your solution to be judged it is **highly recommended** that you perform a test run on your submission first, due to possible differences between the environment you utilised for development and the environment in which your submission will be run on the judges machine.

To test your submission first, perform the following steps:

1. Ensure that your source code file and the sample data files are in the same folder/directory on your system.
2. From the Problem dropdown list select the challenge that you are attempting.
3. From the Language dropdown list select the programming language in the submission is written in.
4. Use the Select button to select the source code file for the submission. (Note: A single Source Code file is required, do not attempt to submit data files or executable files).
5. Use the Add Button to select any additional files needed to complete your submission. (Note: This is rarely needed).
6. Click on "Test".
7. Select the appropriate sample input file in the open dialog box. (Typically the sample input file will be <challenge>_sample_in.txt).
8. Wait for the output dialog and compare to the expected output.
9. If you are happy with your submission, then submit your solution for judging, by clicking on "Submit".

2.2 Software Languages

Currently the competition support the following software development languages with some variations between regional areas: Java, C++ and C#.

2.2.1 Java

Java is compiled utilising the Oracle Java 6SE JRE implementation, however future competitions may migrate to Java 7SE as Java 7 becomes more popular. (This guide will target the Oracle Java 6 SE JRE).

By default the competition will utilise the 32bit JRE, however this may vary as needed between each region. Additionally the Java compiler and JVM are run using default settings only.

2.2.2 C++

C++ (and by extension C) is compiled with an POSIX compatible compiler, typically being mingw on Windows. mingw utilises the GNU GCC compiler suite, and offers a near complete POSIX environment including the C++ STL.

It should be noted, that in some instances the Microsoft Visual Studio C++ compiler has been used within the competition, so it is best to check with the local staff supporting the competition which compiler will be utilised.

Irrespective of the C++ compiler and/or environment, it should be noted that the default compiler settings are utilised through the competition, so features including optimisation flags or 64bit operation are not enabled.

This guide will target a 100% pure POSIX environment.

2.2.3 C#

C# will typically be compiled by Microsoft Visual Studio 2010 with the .NET 2.0 Framework. However there may be variations to this, so it is best to check with the local staff supporting the competition which compiler and/or .NET framework will be utilised.

2.3 IDEs

At the moment there are no official supported IDEs utilised by the competition, however the majority of contestants utilise either Eclipse and/or Visual Studio.

Other IDEs or Editors commonly utilised by competitors include NetBeans (Java, C++), Code::Blocks (C++) and Notepad++ (Java, C#, C++).

2.3.1 Eclipse

Eclipse may be utilised to develop either Java applications or C++ applications (on provision the appropriate eclipse plugins for C++ are installed, and a compatible C++ compiler such as mingw is also installed).

There is no special configuration for Eclipse to be utilised within the competition. As all competition entries operate within a command line only interface there is no requirement for any GUI builder plugins to be present.

To utilise Eclipse for Java development, perform the following steps:

1. Start Eclipse, and switch to a Workspace that is empty, or has been designated for use for competition. (Use File -> Switch Workspace to move).
2. In the File menu, select New.
3. In the New dialog box, select Java Project. Click on Next.
4. Enter any name for the project name. Leave all other settings as default.
5. Click on Finish. This will create a basic project that can be used for developing submissions.
6. In the Package Explorer pane, right click on the 'src' package, and select New -> Java Class. This will be the first submission that you will work on. When developing further submissions, simply start at this point and following the remaining steps.
7. In the New Java Class dialog, enter in the challenge name in the Name: field, and click on "public static void main(String[] args)" to select this option. Leave all other options as default, and click on Finish. (See Figure 2.5).

8. The new submission Java file will open in the file pane. (See Figure 2.5)
9. You are now reading to develop you submission.

All debugging facilities may be utilised within Eclipse with no restrictions.

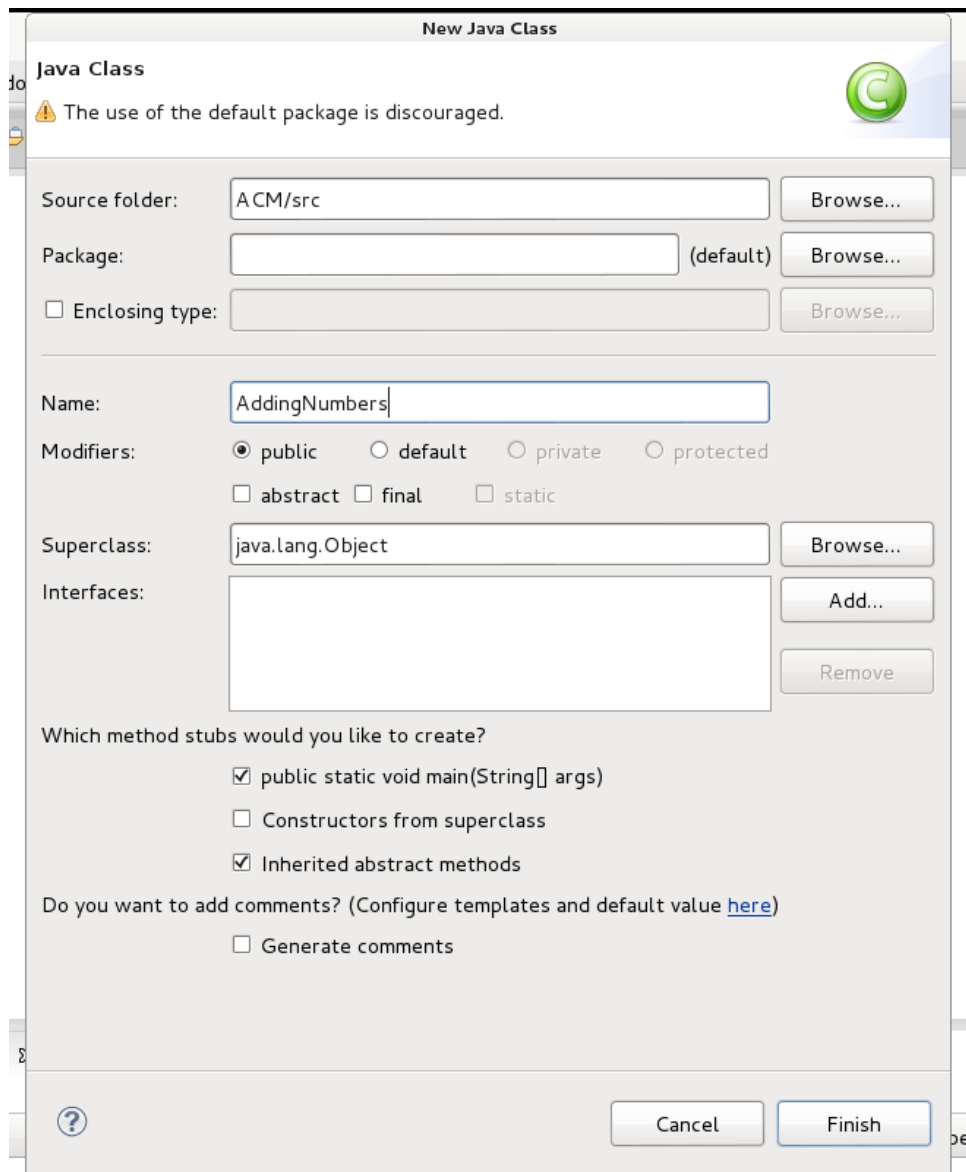


Figure 2.4: Eclipse - New Java Class

One item to note with the file structure of Eclipse and essentially all Java application development, you will need to note the exact location of the source files so are able to find them later in order to submit your solutions.

Using the example in Figure 2.4, note the “Source Folder” location, this is the location that your submissions will be located in, in this case ACM/src/AddingNumbers.java.

When testing your submission with PC² you will be required to either:

1. Copy your source code file to the same location as sample input files, or
2. Copy the sample input files into the same location as the source code file.




Figure 2.5: Eclipse - New Java Application

2.3.2 Visual Studio

Visual Studio is capable of working with a number of programming languages, including C# and C++. This guide will cover working with both languages, noting the differences between the two.

2.3.2.1 C# There is no special configuration for Visual Studio to be utilised within the competition. As all competition entries operate within a command line only interface there is no requirement for any GUI builder plugins to be present.

To utilise Visual Studio for C# development, perform the following steps:

1. Start Visual Studio.
2. If you receive a "Select Development Language" dialog, select Microsoft C#.
3. Select "New Project" from the Start Page, or alternatively from the File menu.
4. In the "New Project" dialog, ensure that:
 - (a) Visual C# > Windows is selected in the Installed Templates pane.
 - (b) Console Application - Visual C# is selected in the Application Type pane.
 - (c) .NET Framework 2.0 is selected in the .NET Framework dropdown.
 - (d) Enter the name of the challenge in the Name: field.
 - (e) Note the location in which the project is being created.
 - (f) Click on OK to build the project. (See Figure 2.6)
5. The new Program.cs file will be opened and displayed in the File Pane. (See Figure 2.7)




Figure 2.6: Visual Studio C# New Project

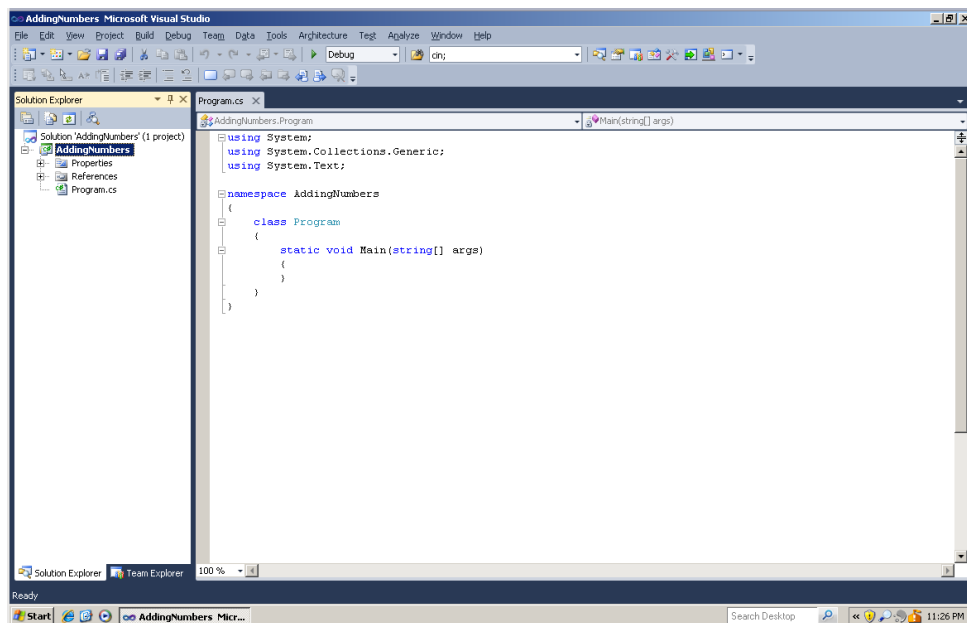


Figure 2.7: Visual Studio C# Program.cs

When testing your submission with PC² you will be required to either:

1. Copy your source code file to the same location as sample input files, or
2. Copy the sample input files into the same location as the source code file.

2.3.2.2 C++ To utilise Visual Studio for C++ development, perform the following steps:

1. Start Visual Studio.
2. If you receive a "Select Development Language" dialog, select Microsoft C++.
3. Select "New Project" from the Start Page, or alternatively from the File menu.

4. In the “New Project” dialog, ensure that:
 - (a) Visual C++ > Win32 is selected in the Installed Templates pane.
 - (b) Win32 Console Application - Visual C++ is selected in the Application Type pane.
 - (c) .NET Framework 2.0 is selected in the .NET Framework dropdown.
 - (d) Enter the name of the challenge in the Name: field.
 - (e) Note the location in which the project is being created.
 - (f) Click on OK to build the project. (See Figure 2.8)
5. The Win32 Application Wizard will run. Select Next on the Wizard Welcome screen.
6. On the Applications Settings dialog, ensure that “Console Application” is checked, and “Precompiled Headers” is unchecked. (See Figure 2.9). Click on Finish to build the project.
7. The new <Application>.cpp file will be opened and displayed in the File Pane. (See Figure 2.10)
8. The following changes are recommended to the main source file to ensure maximum platform compatibility:
 - (a) Add “using namespace std;” before any included headers.
 - (b) Remove or comment out the “#include "stdafx.h"” line
 - (c) Add “#include <stdio.h>” and “#include <iostream>” to ensure POSIX compatibility.
 - (d) Change `int _tmain(int argc, _TCHAR* argv[])` to `int main()`
9. You are now ready to develop your submissions.

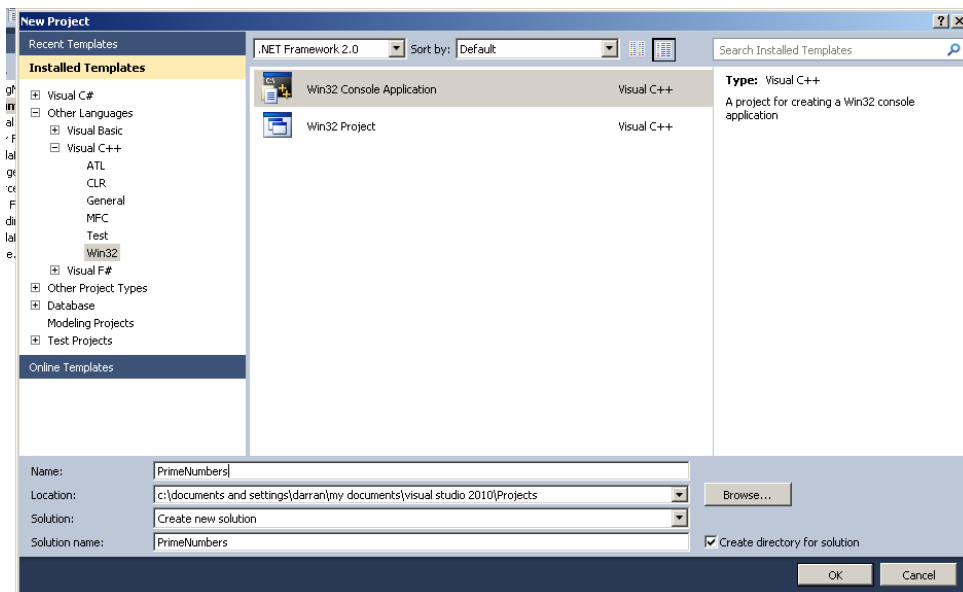


Figure 2.8: Visual Studio C++ New Project




Figure 2.9: Visual Studio C++ Project Options




Figure 2.10: Visual Studio C++ Program.cpp

When testing your submission with PC² you will be required to either:

1. Copy your source code file to the same location as sample input files, or
2. Copy the sample input files into the same location as the source code file.

2.4 Submission Guidelines

The primary requirement of any submission, is that all source code required for the submission is located in a single text file.

To test your submission, the judges machine will compile your source code to an executable or class file in the case of Java, then execute it. All input for the application is feed in via `stdin` (or using standard console input), and all application responses and feedback should be returned via `stdout` (or using standard console output). In effect the judges machine will run:

```
$      submission.exe < challenge1_input.txt > challenge1_output.txt
```

Any output the application produces is saved to a file, and this file is then compared to a known correct answer file. Any variations from between the applications output and the answer file will result in a **No** response. If both the output of the application and answer file match, then the judges machine will return a **Yes** response.

3 Performance

3.1 BigO Notation

Big O notation is used in Computer Science to describe the performance or complexity of an algorithm. Big O specifically describes the **worst-case** scenario, and can be used to describe the execution time required or the space used (e.g. in memory or on disk) by an algorithm.

Typically, algorithms will be assigned the following functions:

O(1) This algorithm operates in the same time irrespective of the number of elements to be processed. (The ideal algorithm).

O(log₂n) This algorithm will have a worst-case runtime of log₂n with n elements to be processed.

O(n) This algorithm will have a worst case runtime in the order of the number of elements to be processed.

O(nlog₂n) This algorithm will have a worst case runtime in the order of nlog₂n with n elements to be processed.

O(n²) This algorithm will have a worst case runtime in the order of n² with n elements to be processed.

O(n³) This algorithm will have a worst case runtime in the order of n³ with n elements to be processed.

O(2ⁿ) This algorithm will have a worst case runtime in the order of 2ⁿ with n elements to be processed.

As an example, suppose the each operation can be done in 1 microsecond, and we have 256 elements to be processed. The respective worst case runtimes for each function is shown in Table 1.

Function	Time
log ₂ n	8 microseconds
n	256 microseconds
nlog ₂ n	2 milliseconds
n ²	65 milliseconds
n ³	17 seconds
2 ⁿ	3.7×10 ⁶⁴ centuries

Table 1: Big O runtimes

The information within this section is a very brief introduction to BigO notation, is included to help you choose between two different algorithms to complete a task based on the Order function given in the algorithm description. Much research has been completed on algorithm analysis, with many text books and reference books being authored on this one subject of Computer Science.

3.2 Measuring Performance

While information described with algorithms can be useful to gain an understanding of the algorithm complexity, it can be helpful to be able to measure the actual execution time needed to complete sections of code. All modern languages or software libraries contains functions to determine execution times².

²Modern IDEs such as Oracle Solaris Studio 12.3 include profiling tools to determine 'hot-spots' within applications and also be able to automatically record execution times of individual functions/methods to later analysis.

3.2.1 C#

The .NET Framework provides a Stopwatch class that is capable of being used to measure execution times. Listing 1 shows an example of the syntax.

```
Stopwatch st = new Stopwatch();
st.Start();
// code to be timed goes here
st.Stop();
// time in milliseconds
long elapsed = (long) st.ElapsedTicks * 1000000 / Stopwatch.Frequency;
Console.WriteLine("time_={0}", elapsed);
```

Listing 1: Timing - C#

3.2.2 C++

The C standard library includes time specific functions in <time.h> on most systems. The primary function is the clock() function that returns the number of 'clicks' since the application started execution. A macro CLOCKS_PER_SEC is used to determine the ratio between clicks and seconds. Listing 2 shows an example of the syntax.

```
#include <time.h>
#include <stdio.h>
clock_t start = clock();
// code to be timed goes here
clock_t end = clock();
// time in clicks
long elapsed = (long)(end - start)/CLOCKS_PER_SEC;
printf("time_=%d\n", elapsed);
```

Listing 2: Timing - C++

3.2.3 Java

The Java System library provides multiple timers with varying accuracy. Since Java 5, nanosecond timers has been available via the System.nanoTime() method³. Listing 3 shows an example of the syntax.

```
long start = System.nanoTime();
// code to be timed goes here
long end = System.nanoTime();
// time in nanoseconds
long elapsed = (end - start)/1000000;
System.out.printf("time_=%d\n", elapsed);
```

Listing 3: Timing - Java

3.2.4 Integers vs Floating Point

Modern CPUs have integrated high-performance floating-point execution units, however it should be noted that the choice of Integer, Long's and Floating Point number will have an impact of performance of your solution. The code snippet in Listing 4 demonstrates the performance differences of using 'int', 'long' and 'double' for a simple add, multiple and divide sequence.

³Accuracy of the System.nanoTime() method is reliant on the JVM version and underlying Operating System. However most modern operating systems do provide some form on nanosecond timer.

```
public class NumberTypeTesting {
    static final long ITERATIONS = 1000000000;
    public static void main(String[] args) {
        final double DOUBLE_PRIME = 73;
        final int INT_PRIME = 73;
        final long LONG_PRIME = 73;
        long start;
        long end;
        long count;
        double valueDoubleA = 1.00;
        double valueDoubleB = Math.PI; // pi = 3.142
        double valueDoubleC = Math.E; // e = 2.718
        int valueIntA = 1;
        int valueIntB = 10;
        int valueIntC = 31;
        long valueLongA = 1;
        long valueLongB = 31;
        long valueLongC = 33;

        // Integers
        count = ITERATIONS;
        start = System.nanoTime();
        while (count-- != 0) {
            valueIntA += valueIntC * valueIntB / INT_PRIME;
        }
        end = System.nanoTime();
        System.out.printf("Integer time = %d msec\n", (end - start) / 1000000);

        // Long
        count = ITERATIONS;
        start = System.nanoTime();
        while (count-- != 0) {
            valueLongA += valueLongC * valueLongB / LONG_PRIME;
        }
        end = System.nanoTime();
        System.out.printf("Long time = %d msec\n", (end - start) / 1000000);

        // Double
        count = ITERATIONS;
        start = System.nanoTime();
        while (count-- != 0) {
            valueDoubleA += valueDoubleC * valueDoubleB / DOUBLE_PRIME;
        }
        end = System.nanoTime();
        System.out.printf("Double time = %d msec\n", (end - start) / 1000000);
    }
}
```

Listing 4: Source code for Timing test

The results⁴ for the above test in Listing 4 are:

```
Integer time = 1131 msec
Long time = 2673 msec
Double time = 1298 msec
```

While there is a minor performance drop for using floating point numbers, using 64bit longs yields over double the execution time. This may easily be fixed by utilising an environment that runs as 64bit code, but this is not guaranteed to be available during the competition.

⁴Java 6SE 32bit was used to generate the following results.

3.3 Implementation and Modern Software Engineering Practices

One of the aims of the competition is to develop efficient solutions to the challenges being presented. However often this also means not following modern software engineering practices and taking as many shortcuts as possible.

Some items that are typically seen (and encouraged) are:

1. Liberal use of global variables utilised by direct access.
2. Libreal use of function pointers and jump tables in C++.
3. Dispite strong OOP principles with each programming language, these are often ignored for more simple data structures and items like inheritance and encapsulation are ignored.
4. Nested classes liberally use public variables allowing for direct access.
5. The “goto” statement being used in C++ and C# submissions⁵.
6. Ignoring typical design patterns, unless they provide direct and significant benefit in utlising an algorithm to complete a challenge.
7. Error checking is kept to a minimum, mainly designed around corner cases for algorithms to handle rather than handling bad and malformed input.
8. Utilisation of the most efficient algorithm, even in cases where memory requirements may be pushed to the extreme. eg, building complete hashtables in memory for $O(1)$ lookups, rather than recomputing as needed.

As mentioned in Section 2.4, that all source code required for the submission is located in a single text file, also requires some creative uses of both local and anonymous classes.

⁵While most professional programmers avoid “goto” as it’s considered inherently evil, there are some instances where its use can save execution time and/or reduce code complexity.

4 Basic Source Templates

All source code submissions are to consist of a single source code file, as previously mentioned. This section aims to provide simple templates that can be utilised to create your submissions. It will also cover some of the basic console functions available with each language.

4.1 Input / Output

4.1.1 C#

The .NET Framework unfortunately has rather cumbersome support for handling console input and output. The `System.Console`⁶ class provides methods for dealing with the console. The three main methods that are typically used are:

1. `Console.ReadLine();`
2. `Console.Write();`
3. `Console.WriteLine();`

4.1.1.1 Input The primary function for input from the Console is the `Console.ReadLine()` method which as the name indicates, reads a single line from the console and returns a string.

In order to extract information from the string, it is needed to split the string based on a delimiter (typically a space), then attempt to convert each part into the desired type. Listing 5 shows how to read a group of 3 integers (per line) from the console, until a three 0's (zeroes) are entered. The numbers for each line is added, and the sum is written back to the console.

4.1.1.2 Output The primary functions for output to the Console are the `Console.Write()` and `Console.WriteLine()` methods. These two differ only by the latter terminating the line with a carriage-return, while the former does not.

One item to note, that a single `Console.Write()` method may only take up to 5 parameters, the first being a string, and the other 4 being items to be inserted into the string. Item placement within the string is denoted by a number with `{}` brackets. (See the last line in Listing 5 for an example). The item placement parameter, may also take a second argument, being the type to have the item converted to, or displayed as. Common types include:

```
(C) Currency: . . . . . {0:C}
(D) Decimal: . . . . . {0:D}
(E) Scientific: . . . . . {0:E}
(F) Fixed point: . . . . . {0:F}
(G) General: . . . . . {0:G}
(P) Percent: . . . . . {0:P}
(X) Hexadecimal: . . . . . {0:X}
```

By default, console output is buffered, and only written periodically as determined by underlying system settings. To flush the output to console immediately, the `Console.Out.Flush()` method can be utilised.

⁶<http://msdn.microsoft.com/en-us/library/system.console.aspx>


```
using System;
namespace AddNumbers {
    class Program {
        static void Main(string[] args) {

            // Define our numbers to read.
            int[] numbers;
            int index;
            int sum;
            string line;
            string[] linesplit;

            // Keep reading the input from the console until we have nothing left.
            while ((line = Console.ReadLine()) != null) {

                // Test for exit condition.
                if (line.CompareTo("0_0_0") == 0) {
                    break;
                }

                // Reset array indices and sum of numbers
                index = 0;
                sum = 0;

                // Split the line read, and create a new int array to hold our value.
                linesplit = line.Split(' ');
                numbers = new int[linesplit.Length];

                // Attempt to convert the individual parts to int's
                foreach (string element in linesplit) {
                    try {
                        numbers[index] = Convert.ToInt32(element);
                    }
                    catch {
                        numbers[index] = 0;
                    }
                    index++;
                }

                // Sum our number and output the sum to Console.
                foreach (int number in numbers) {
                    sum += number;
                }
                Console.WriteLine("{0}", sum);
            }
        }
    }
}
```

Listing 5: C# Input Example

4.1.2 C++

Due to the environment in which C++ was originally developed, C++ has very strong capabilities for handling both console input and output. C++ offers two methods when working with the console:

1. `iostreams`
2. C standard library functions.

While the two methods can be intermixed, it is recommended that programmers utilise a single method for their application⁷. For the purposes of this guide, I'll only explain the `iostreams` method as it is often seen as easy to use of the two methods.

4.1.2.1 Input Input from the console is handled by the `std::cin` stream, and has the ability to take multiple types of inputs in a single line or function call. (This is possible due to operator overloading in C++). Listing 6 shows the same application written in C++, as shown in C# within Listing 5.

```
,
#include <cstdlib>
#include <iostream>
using namespace std;

int main() {
    int a, b, c;

    do {
        // Get input of 3 integers and store in a, b and c.
        cin >> a >> b >> c;

        // Test exit condition, and exit if true.
        if(a == 0 && b == 0 && c == 0)
            break;

        // Output the sum.
        cout << (a + b + c) << endl;
    } while(true);
    return 0;
}
```

Listing 6: C++ Input and Output Example

Since the `cin` and `cout` streams operate on single variables, an alternate method is required to read a complete line in one function call. This method is `getline (istream& is, string& str);`, where `is` is the character stream, and `str` is the string to place the input into. As example of `getline()` is displayed in Listing 7.

```
,
#include <iostream>
#include <string>
using namespace std;

int main () {
    string str;
    cout << "Please enter full name: ";
    getline (cin, str);
    cout << "Thank you, " << str << ".\n";
    return 0;
}
```

Listing 7: C++ `getline()` example

⁷Mixing the two methods is possible, on provision that all input and output streams are empty when switching between either method. This is due to the buffering that each method utilises during Console IO operations.

Once the line has been fetched with the `getline()` method, you are free to use any of the other string functions to extract information from the string.

4.1.2.2 Output The primary method of output is via the `cout` iostream, as shown in Listing 6. The format is simply:

```
cout << {object} << {object} << {object} << " string " << .. << endl;
```

where each `{object}` represents any C++ primitive type or any C++ object.

There are two methods to generating a newline character, utilising C++. You may either:

1. Output string `"\n"`, or
2. Output `std::endl`.

The difference between the two methods, is that `"std::endl"` will flush the output buffer to console, where `"\n"` will not.

To control the precision of floating point numbers, you can use the `setprecision(x)` method as part of the output sequence. eg: `cout << setprecision(4) << (double)1.23456788 << endl;` will output 1.235 to the console.

4.1.2.3 iostreams vs printf The other method to perform console output is the C function `printf()`. `printf()` offers the same features as the `cout` iostream, and may be used when very fine control over output is required especially with floating point numbers.

The general format of the `printf()` function is:

```
printf(const char *str, ...);
```

`str` is a formatted string, that may contain 0 or more place holders for additional arguments. Placeholders in the formatted string are simply filled in order of additional arguments as specified in the function call, and the additional arguments must be of the same type as specified by the placeholder.

Formats for placeholders include:

```
%d - decimal
%du - decimal unsigned
%f - floating point
%s - string (char*)
%c - character
%x - hexadecimal number
%l - long
%lu - long unsigned
```

Additional fields may be added to the place holders to specify field width and/or precisions. For example:

```
%.5f - will display a floating point number to 5 decimal places.
%5s - will consume at exactly 5 character spaces for a string.
```

There are a few special reserved characters for the `printf()`, some of these include:

```
\n - carriage return.
\t - tab character.
```

There has been some debate over the performance aspects of the two methods, often citing that there is no performance difference. Utilising the source code in Listing 8, I've found that in some cases there can be significant differences in performance, with `printf()` being up to 250 times faster than the equivalent `cout` function.

Based on average times for 5 runs of the test application in Listing 8, yields the following results:

1. cout time = 67240 msec
2. printf() time = 280 msec

```
#include <cstdlib>
#include <iostream>

using namespace std;

#define LOOPCOUNT 1000000
#define NUMBER1 31.0
#define NUMBER2 21

int main() {
    long count = LOOPCOUNT;
    double dbl = NUMBER1;
    int Int = NUMBER2;
    char* str = (char*) &"hello";

    clock_t start = clock(); // Start timing
    while (count--) {
        cout << dbl << " " << Int++ << " " << str << "\n";
        dbl = dbl * (double) Int / ((double) Int * 2.0);
    }
    clock_t end = clock(); // End timing
    cerr << "cout_time=" << double(end - start)*1000.0/CLOCKS_PER_SEC
        << "msec" << endl;

    // reset start values
    count = LOOPCOUNT;
    dbl = NUMBER1;
    Int = NUMBER2;
    start = clock(); // Start timing
    while (count--) {
        printf("%f%d%s\n", dbl, Int++, str);
        dbl = dbl * (double) Int / ((double) Int * 2.0);
    }
    end = clock(); // End timing
    cerr << "printf_time=" << double(end - start)*1000.0/CLOCKS_PER_SEC
        << "msec" << endl;

    return 0;
}
```

Listing 8: C++, cout vs printf()

4.1.3 Java

Similar to C++, Java has a very capable set of support functions for handling console input and output. These are mainly archived through the `System.in` and `System.out` classes used in conjunction with the `Java.util.Scanner` class provided with the default Java libraries.

4.1.3.1 Input Java historically has had a large number of different methods for handling console input, with each new version of Java providing a more streamlined method of handling these functions.

The current preferred method for console input in Java is to use the `Java.util.Scanner` class tied with the `System.in` object to extract the required information from the console. An example of the `Scanner` class can be found in Listing 9 which solves the same problem as shown in Listing 5.

```
import java.util.Scanner;
public class AddNumbers {
    public static void main(String[] args) {

        Scanner in = new Scanner(System.in);

        // get first line and check for end of test cases
        String line = in.nextLine();

        // Continue until exit condition
        while (!line.equals("0_0_0")) {

            // extract three ints
            Scanner sc = new Scanner(line);
            int a = sc.nextInt();
            int b = sc.nextInt();
            int c = sc.nextInt();

            System.out.printf("%d\n", (a + b + c));

            // get next line
            line = in.nextLine();
        }
    }
}
```

Listing 9: Java Input and Output Example

4.1.3.2 Output Output is easiest handled via the `System.out.printf()` method, as it offers a good match between flexibility and performance. The format for the method call is the same as the C++ `printf()` function as described in Section 4.1.2.3. An example of the method call is also in Listing 9.

5 Basic Algorithms

5.1 Sorting

The majority of data structures that will be utilised within the competition all provide some form of inbuilt sorting algorithm, or through their design are naturally sorted as in the case of a Binary Search Tree.

It is highly recommended as far as competition submissions are concerned, that you utilise the built-in sort methods rather than attempting to implement your own sort method.

Typically, most data structures will utilise either quicksort or merge sort (depending on the underlying structure) as they both offer $O(n \log n)$ performance in the average case.

5.2 Searching

When given a linear array of data items, search algorithms find information about a particular data item in the list or find the location of the data item in the list. Two primary search algorithms are:

1. Linear Search
2. Binary Search

5.2.1 Linear Search

Linear Search algorithms transverse through a list of data items in sequential order attempting to find the location of the data item. The list itself may or may not be sorted, and the underlying data structure may be a linear array or a linked list.

Algorithm 1 Linear Search

Input Vector S , with n elements, with search key k

Output if $k \in S$ return index of k , else return -1

```
procedure LINEARSEARCH( $S, k$ )
  for  $i = 0$  to  $n - 1$  do
    if  $S[i] = k$  then
      return  $i$ 
    end if
  end for
  return -1
end procedure
```

5.2.1.1 Description of working The linear search algorithm takes a vector (aka array) of elements, and simply searches all elements in order as stored. This can yield slow performance with large vectors, as the worst case for linear search is $O(n)$.

5.2.1.2 Implementation Listing 11 and Listing 11 show the linear search algorithm as implemented in Java and C++ respectively.

```
public int LinearSearch(E[] vector, E key){
    for(int index = 0; index < vector.length; index++){
        if(vector[index] == key){
            return index;
        }
    }
    return -1;
}
```

Listing 10: Linear Search Implementation (Java)

```
public int LinearSearch(E[] vector, int vsize, E key){
    for(int index = 0; index < vsize; index++){
        if(vector[index] == key){
            return index;
        }
    }
    return -1;
}
```

Listing 11: Linear Search Implementation (C++)

5.2.1.3 Sample Problem - Linear Search Given a list of numbers (integers) in a line, determine if the first value on the line is present within the subsequent list of numbers.

INPUT

Input consists of one or more lines, with the first line being the number of cases to test.

Each line consists of one or more integers in the range of 0 to 32767. The first integer is the key value, followed by a list of up to 32 integers forming a vector of numbers.

SAMPLE INPUT

```
3
10 12 327 0 10
1 2 3 4 5 6 7 8
10 20 30 40 50 60 70 90 10
```

OUTPUT

The output of each line should consist of a single integer being either the index of the key within the vector, being zero (0) offset, or the value -1 if the key is not present in the vector.

SAMPLE OUTPUT

```
3
-1
7
```

5.2.1.4 Problem Solution A solution to the above problem utilising a linear search can be seen in Listing 12. The solution included simply reads in a line of numbers, and attempts to find the first value in the list of other values in the line. It continues to do this, until the exit condition is reached.

Some of the test cases to handle include:

1. The case count being less or equal to 0.
2. The line itself contains a single integer, being the key, but provided with an empty vector to search.

```
import java.util.Scanner;

public class LinearSearch {

    /**
     * Perform linear search of array (vector) for item (key).
     *
     * @param vector array of numbers
     * @param key item to look for in array
     * @return index of key in vector, or -1 is not present
     */
    public static int LinearSearch(int[] vector, int key) {
        for (int index = 0; index < vector.length; index++) {
            if (vector[index] == key) {
                return index;
            }
        }
        return -1;
    }

    /**
     * Main
     */
    public static void main(String[] args) {
        Scanner in = new Scanner(System.in);

        // get first line and get the number of cases to test.
        int caseCount = Integer.parseInt(in.nextLine());

        // Keep reading each line while caseCount > 0
        while (caseCount-- > 0) {

            // split by white space. so we have an array of numbers
            String[] numStrs = in.nextLine().split("\\s+");

            // create an array to hold our numbers, and convert the array of strings to
            // numbers. Note: numStrs[0] is the key value
            int[] nums = new int[numStrs.length - 1];
            for (int i = 1; i < numStrs.length; i++) {
                nums[i - 1] = Integer.parseInt(numStrs[i]);
            }

            // Output the index of the key in the vector
            System.out.printf("%d\n", LinearSearch(nums, Integer.parseInt(numStrs[0])));
        }
    }
}
```

Listing 12: Solution to Linear Search (Java)

5.2.2 Binary Search

The binary search algorithm is a more efficient method of searching a vector, on provision that the vector is sorted and any element can be accessed in $O(1)$ time. Because of these two conditions, it can't work with some storage data structures like linked lists (as elements can't be accessed in $O(1)$ time), nor is suitable for vectors that are unsorted.

Algorithm 2 Binary Search

Input An ordered vector S , with n elements, with search key k . Items low and $high$ indicate current search space of vector S

Output if $k \in S$ return index of k , else return -1

```
procedure BINARYSEARCH( $S, k, low, high$ )
  if  $low > high$  then
    return -1
  else
     $mid \leftarrow \lfloor (low + high) / 2 \rfloor$ 
    if  $k = S[mid]$  then
      return  $mid$ 
    else if  $k < S[mid]$  then
      return BINARYSEARCH( $S, k, low, mid - 1$ )
    else
      return BINARYSEARCH( $S, k, mid + 1, high$ )
    end if
  end if
end procedure
```

5.2.2.1 Description of working The binary search algorithm is a naturally recursive algorithm, in that it calls itself to continue searching the vector.

The algorithm starts with the entire space of the vector, and looks at the mid point between the low and $high$ values. If this value is not the required key (k), it will then determine if the key is less than or greater than the current value at mid . If the key is lower than mid , then it will redefine the search space to be that between low and $mid-1$, otherwise redefine the search space between $mid+1$ and $high$. It then calls itself to perform another search. In the event that low is greater than $high$, it determines that the key is not in S , and will return -1.

What the algorithm effectively does is split the entire search space of vector S into 2 parts, if the key is not at mid . By virtue, if the value at mid is less than the key, it understands that there is no justification to look at values located to the left of the current mid point in the vector. With each iteration of the search it effectively reduces the search space by half.

By reduction of the search by half, the worst case performance of a binary search is $O(\log n)$. As this is a vast improvement on a linear search, a binary search should be utilised when ever possible. However, this requires that the vector be sorted before a binary search can be performed.

Utilising a quicksort or merge sort, will add overhead (both of these typically yield $O(n \log n)$ performance), so for very large vectors, the overhead of a sort prior to search is not that great, but for small size vectors, the overhead of a sort may not yield greater performance over the simple (and slow) linear search.

5.2.2.2 Implementation Listing 13 shows an implementation of the binary search in Java.

```
public int BinarySearch(E[] vector, E key, int low, int high){
  if(low > high){
    return -1;
  }
  int mid = (low+high)/2;
  if(vector[mid] == key){
    return mid;
  } else {
    if(key < vector[mid]){
      return BinarySearch(vector, key, low, (mid-1));
    } else {
      return BinarySearch(vector, key, (mid+1), high);
    }
  }
}
```

```
}  
}
```

Listing 13: Binary Search (Java)

5.2.2.3 Sample Problem For demonstration of the Binary Search algorithm, I will use the same problem as shown in Section 5.2.1.3, Linear Search Problem. However, it is expected that the input of numbers (except for the key) will be in order from lowest to highest.

5.2.2.4 Problem Solution A solution to the sample problem utilising a binary search can be seen in Listing 14. The solution included simply reads in a line of numbers, and attempts to find the first value in the list of other values in the line. It continues to do this, until the exit condition is reached.

Some of the test cases to handle include:

1. The case count being less or equal to 0.
2. The line itself contains a single integer, being the key, but provided with an empty vector to search.

```
import java.util.Scanner;  
  
public class LinearSearch {  
  
    /**  
     * Perform binary search of array (vector) for item (key).  
     *  
     * @param vector array of numbers  
     * @param key item to look for in array  
     * @param low start position of array to search  
     * @param high end position of array to search  
     * @return index of key in vector, or -1 is not present  
     */  
    public static int BinarySearch(int[] vector, int key, int low, int high){  
        if(low > high){  
            return -1;  
        }  
        int mid = (low+high)/2;  
        if(vector[mid] == key){  
            return mid;  
        } else {  
            if(key < vector[mid]){  
                return BinarySearch(vector, key, low, (mid-1));  
            } else {  
                return BinarySearch(vector, key, (mid+1), high);  
            }  
        }  
    }  
}  
  
/**  
 * Main  
 */  
public static void main(String[] args) {  
    Scanner in = new Scanner(System.in);  
  
    // get first line and get the number of cases to test.  
    int caseCount = Integer.parseInt(in.nextLine());  
  
    // Keep reading each line while caseCount > 0  
    while (caseCount-- > 0) {
```

```
// split by white space. so we have an array of numbers
String[] numStrs = in.nextLine().split("\\s+");

// create an array to hold our numbers, and convert the array of strings to
// numbers. Note: numStrs[0] is the key value
int[] nums = new int[numStrs.length - 1];
for (int i = 1; i < numStrs.length; i++) {
    nums[i - 1] = Integer.parseInt(numStrs[i]);
}

// Output the index of the key in the vector
System.out.printf("%d\\n",
    BinarySearch(nums, Integer.parseInt(numStrs[0]), 0, nums.length-1));
}
}
```

Listing 14: Solution to Binary Search Problem (Java)

5.3 Array Handling

When talking about arrays, we typically define one as a single string⁸ or allocation of elements in a linear continuous region. We can define an array of arrays to form a two dimensional array, or an array of arrays of arrays to form a three dimensional array, and continue to do so, allowing for infinite dimensional arrays. This section will typically discuss array operations on two dimensional arrays, such as rotation and mirroring that may support application of algorithms or may simply speed up implementations due to underlying hardware constraints.

5.3.1 Array Performance

When most people are taught programming in either High School or early University level, performance constraints in regards to arrays is either neglected or very limited discussion is made without concrete examples. This section aims to give some insight to performance issues when dealing with arrays, primarily around performance bottlenecks.

The primary reason for poor performance when using arrays, is not based on a programming language or library issue, but is based on lack of understanding the underlying hardware and how memory access works.

From a hardware architectural viewpoint there are different classes of hardware memory:

1. Primary - The RAM that the CPU sees as the address space given to it.
2. Secondary - The Harddrive installed within the system, providing non-volatile memory.
3. Tertiary - Removable non-volatile memory such as DVD's, CD's, USB Flash Keys, etc.

What is typically not taught are the different levels of primary storage. Utilising a modern Intel x86 processor⁹ as an example, the primary levels include:

1. CPU registers, internal to the CPU and these are where typically most operations are performed. These typically have a zero latency access.

⁸A string in many languages (notably low level languages like assembler) does not define a String of letters, numbers and punctuation, but rather a linear memory region of bytes.

⁹The size and access latencies described are taken from the Intel Core 2 Quad Family Datasheet and Vol 3 of the Intel Architecture Manuals.

2. Level 1 Data and Code caches, these hold the most recent code and data being accessed from the Level 2 cache. There are typically two Level 1 caches, each designated for holding either code or data, and are typically 16kilobytes size in size¹⁰. The Level 1 cache will typically have a 1-2 CPU cycle access latency, and besides the CPU registers is the fastest memory available to the CPU.
3. Level 2 unified Code and Data cache will typically be in the size of anywhere from 256kilobytes up to 16megabytes in size depending on the CPU make and model. This acts a large cache between the main memory of the system, and the CPU and Level 1 caches. The Level 2 cache will typically have a 5-10 CPU cycle access latency, primarily due to restrictions of the size of the cache. (The larger the cache, the slower the access due to it's size¹¹).
4. Level 3 unified Code and Data caches are present in some CPUs and act as a third level between the CPU and main memory of the system. Level 3 caches are becoming more common with multi-core CPUs, as Level 2 and Level 1 caches are being tied to a particular CPU core, where the Level 3 cache can act as a unified cache for all CPU cores. While Level 3 caches can be quite big, in some cases now approaching 32MB in size¹², they are even slower that the Level 2 cache with access latency between 15-35 CPU cycles.
5. Main Memory, is typically in the form of the DIMMs that get installed in the mainboard of the computer. While systems are approaching very large capacities (32GB can be found in home desktop systems, and up to 194GB in workstations), they are very slow compared to the CPU registers and even the Level 1,2 and 3 caches. Typical access latencies can be measured anywhere from hundreds to thousands of CPU cycles. That is, if the CPU needs some information that is not present in one of the caches, it can potentially stall for 100's, if not 1000's of CPU cycles doing nothing while it waits for the information¹³.

Why is knowing all about the different level of cache's important in regards to arrays? Simply, if you try to access an array element that is not in one of the caches, your application will suffer a performance hit whilst waiting for the information from main memory. By ensuring that your next memory access will be in one of the CPUs cache, you can ensure the best possible performance for your applicaton when dealing with any sized array.

To quote Terje Mathisen (a well known programming optimization guru¹⁴): "All programming is an exercise in caching."

The problem with current programming languages, notably Java and C# is that they run on top of virtual machines or utilise some form of JIT compilation, negating any direct control of the CPU and cache management functions. Even C++ applications lack cache management functions (unless you utilise inline assembler in your application). The way to work with these languages is to exploit the nature of the CPU's cache management engine to your advantage.

The CPUs cache management engine works by loading the contents of most recently accessed memory address into the cache in either 32byte chunks for the level 1 cache, or 4kilobyte chunks for the level 2 and 3 caches. To ensure that the next memory address is located in the cache, ensure that the next array element to access is located very close of the last one accessed. When the CPU cache management engine sees your last access was on a border of a chunk it will load in the next chunk in a linear fashion based on the last accessed memory address.

Therefore the to gain the best possible performance when dealing with arrays, either utilise very small arrays that will fit into the level 1 cache, or only access arrays in a linear fashion row by row.

To illustrate these cache performance aspects, the application in Listing 15 yeilds the results show in Table 2. As can be seen, when dealing with a 16384×16384 sized array (consuming 256MB), accessing the array row by row takes 1.35 seconds, however accessing it column by column takes just over 16 seconds.

¹⁰Each CPU make/model can have different L1 cache sizes, for example the Intel E7 Xeon CPUs have 32KB for code and 32KB for data in it's L1 cache, and the AMD Opteron utilises 64KB L1 caches.

¹¹This is a very crude approximation, as there are many factors that determine the performance of the L2 cache.

¹²The Sun UltraSPARC IV+ utilises a 32MB L3 cache and the Intel Itanium 9300 utilises a 24MB L3 cache.

¹³CPU vendors do a lot to avoid this, and even resort to techniques such as SMT (aka HyperThreading), or even offer CPU instructions that allow applications to preload the caches with information to avoid these stalls.

¹⁴Terje, at one time worked for iD Software on the original Doom and Quake games and was able to get Quake running a full 3D environment utilising a software based graphics renderer on hardware such as the Intel Pentium 60. (The Intel Pentium 60, ran at 60MHz, roughly 60-90 times slower than current CPUs).

This clearly demonstrates the caches hits/misses taking place and confirms the latencies expected by the cache misses.

```
using namespace std;
#include <time.h>
#include <stdio.h>
#include <iostream>
char array1[1024][1204];
char array2[2048][2048];
char array3[4096][4096];
char array4[8192][8192];
char array5[16384][16384];

int main() {

    // Row by row.
    clock_t start = clock();
    for (int y = 0; y < 1024; y++) {
        for (int x = 0; x < 1023; x++) {
            array1[y][x] = array1[y][x + 1];
        }
    }
    clock_t end = clock();
    long elapsed = (long) (end - start) / (CLOCKS_PER_SEC / 1000);
    printf("1MB_row_time=%dmsec\n", elapsed);
    // Column by Column
    start = clock();
    for (int x = 0; x < 1024; x++) {
        for (int y = 0; y < 1023; y++) {
            array1[y][x] = array1[y + 1][x];
        }
    }
    end = clock();
    elapsed = (long) (end - start) / (CLOCKS_PER_SEC / 1000);
    printf("1MB_col_time=%dmsec\n", elapsed);

    // .... <snip 4MB, 16MB and 64MB loops

    // 256MB Row by row.
    start = clock();
    for (int y = 0; y < 16384; y++) {
        for (int x = 0; x < 16383; x++) {
            array5[y][x] = array5[y][x + 1];
        }
    }
    end = clock();
    elapsed = (long) (end - start) / (CLOCKS_PER_SEC / 1000);
    printf("256MB_row_time=%dmsec\n", elapsed);
    // Column by Column
    start = clock();
    for (int x = 0; x < 16384; x++) {
        for (int y = 0; y < 16383; y++) {
            array5[y][x] = array5[y + 1][x];
        }
    }
    end = clock();
    elapsed = (long) (end - start) / (CLOCKS_PER_SEC / 1000);
    printf("256MB_col_time=%dmsec\n", elapsed);
    return 0;
}
```

Listing 15: Array Access Performance (C++)

	Row x Row	Column x Column
1024×1024 (1MB)	~0 msec	10 msec
2048×2048 (4MB)	20 msec	220 msec
4096×4096 (16MB)	90 msec	830 msec
8192×8192 (64MB)	350 msec	4150 msec
16384×16384 (256MB)	1350 msec	16080 msec

Table 2: Array Access Performance

5.3.2 Array Traversal Methods

Any 2-dimension array, can be accessing in a variety of ways, including column-wise, row-wise and starting from the top to bottom, bottom to top, left to right and right to left.

As seen in Listing 15, row-wise and column-wise access methods where undertaken in a top to bottom, left to right fashion. Formalised algorithms for row-wise and column-wise access are shown in Algorithms 3 and 4 respectively.

To change from a left to right, to a right to left access pattern, simply count down from the width value to 0 for columns. To change from top to bottom, to a bottom to top access pattern, simply count down from the height value to 0 for rows.

The same applies to any 2+ dimensional array, to change the direction of travel, either change from counting from 0 to width/height to counting down from the width/height to 0, and vice versa.

Algorithm 3 Row-wise Traversal of an Array

Input A source matrix S , with n by m elements.

Output Prints value in element of matrix

```
procedure MATRIXROWWISETRAVERSAL( $S, n, m$ )
  for  $row = 0$  to  $m - 1$  do
    for  $column = 0$  to  $n - 1$  do
      print  $S[row][column]$ 
    end for
  end for
end procedure
```

Algorithm 4 Column-wise Traversal of an Array

Input A source matrix S , with n by m elements.

Output Prints value in element of matrix

```
procedure MATRIXCOLUMNWISETRAVERSAL( $S, n, m$ )
  for  $column = 0$  to  $n - 1$  do
    for  $row = 0$  to  $m - 1$  do
      print  $S[row][column]$ 
    end for
  end for
end procedure
```

5.3.3 Diagonal Traversal of an Array

Diagonal traversal of an array is used for many areas including image analysis, map scanning, simple path finding methods¹⁵.

The one issue with this form of traversal through an array, are the underlying performance penalties that will occur due to cache misses, as each subsequent access to the array is non-linear. Like other array traversal techniques, diagonal traverse is an $O(n^2)$ operation.

Algorithm 5 Diagonal Traversal of an Array

Input A source matrix S , with n by m elements.

Output Prints value in element of matrix

```

procedure MATRIXDIAGONALTRAVERSAL( $S, n, m$ )
   $x \leftarrow 0$ 
   $y \leftarrow 0$ 
  while True do
     $v \leftarrow x$ 
     $w \leftarrow y$ 
    while  $v \geq 0$  and  $w < m$  do
      print  $S[w][v]$                                 ▷  $S[\text{row}][\text{column}]$ 
       $v \leftarrow v - 1$ 
       $w \leftarrow w + 1$ 
    end while
    if  $x < n - 1$  then
       $x \leftarrow x + 1$ 
    else if  $y < m - 1$  then
       $y \leftarrow y + 1$ 
    else
      return
    end if
  end while
end procedure

```

5.3.3.1 Description of working As described in Algorithm 5, the algorithm utilises two loops, with the first loop (`while(true)`) determines the start position of the traverse, and the second while loop traverses the actual slice, starting from the top and moving downwards and to the left, (as denoted by the $v \leftarrow v - 1$ and $w \leftarrow w + 1$ operations).

The if-else-if-else statements recalculate the new start position for the next traversal, until the new start position exceeds the bounds of the array, in which case the algorithm exits.

The algorithm as described, starts in the top left corner, and moves towards the bottom right of the array. To modify the algorithm to scan from other origins to opposite corner, the line to be altered is the access function that prints the current element in the array. The start at the various origins and traverse to the opposite corner, the following forms are needed for the print function:

- Top-left origin to bottom-right traverse: print $S[w][v]$
- Bottom-left origin to top-right traverse: print $S[m - w - 1][v]$
- Top-right origin to bottom-left traverse: print $S[w][n - v - 1]$
- Bottom-right origin to top-left traverse: print $S[m - w - 1][n - v - 1]$

Figure 5.1 demonstrates the top-left origin to bottom-right traversal on a 3×4 array, with each slice (or starting location) shown.

¹⁵There are more comprehensive path finding techniques that utilise graphs/networks, so these methods are not commonly used.

The order of traversal for the other origins would result in the following sequences based on the forms listed above:

- Top-left origin to bottom-right traverse: a, b, d, c, e, g, f, h, j, i, k, l
- Bottom-left origin to top-right traverse: j, k, g, l, h, d, i, e, a, f, b, c
- Top-right origin to bottom-left traverse: c, b, f, a, e, i, d, h, l, g, k, j
- Bottom-right origin to top-left traverse: l, k, i, j, h, f, g, e, c, d, b, a




Figure 5.1: Diagonal Traversal of an Array with a Top-Left Origin

5.3.3.2 Implementation Listing 16 demonstrates the implementation of the algorithm as described.

```
public static void ArrayDiagonalTraverse(int[][] matrix, int width, int height){
    int start_x = 0;
    int start_y = 0;
    int column;
    int row;

    // Keep looping until exit condition.
    while (true) {

        // Initialise the starting location in the matrix for the current slice.
        column = start_x;
        row = start_y;

        // Traverse the current slice.
        while (column >= 0 && row < height) {
            System.out.printf("%d_", matrix[row][column]);
            column--;
            row++;
        }

        // Update the start location for the next slice.
    }
}
```



```
    if (start_x < width - 1) {
        start_x++;
    } else if (start_y < height - 1) {
        start_y++;
    } else {
        // Exit the method, as start locations are now out of matrix bounds.
        break;
    }
}
```

Listing 16: Diagonal Traversal of an Array (Java)

5.3.4 Array Rotation

As shown in Section 5.3.1, the way that an array is accessed can effect the performance of your application. Methods to ensure that you can access an array row by row may require rotation of an array, or alternatively rotation of an array may be needed to utilise an algorithm or math function.

Any array rotation or mirroring function requires two copies of the array to be present at one time, the source array and target array. This must be considered in relation to the amount of memory required, as you effectively need double the memory requirement for either operation.

Both rotation and mirroring functions are $O(n^2)$ operations, due to all array members must be accessed to complete the operation successfully.

Algorithm 6 Array Rotation Clockwise

Input A source matrix S , with n by m elements, with target matrix T , with m by n elements.

Output Matrix T represents matrix S rotated 90°clockwise.

```
procedure ROTATEMATRIXCLOCKWISE( $S, T, n, m$ )
     $f \leftarrow m - 1$ 
    for  $y = 0$  to  $m - 1$  do
        for  $x = 0$  to  $n - 1$  do
             $T[x][(f - y)] \leftarrow S[y][x]$  ▷  $S[\text{row}][\text{column}]$ 
        end for
    end for
end procedure
```

Algorithm 7 Array Rotation AntiClockwise

Input A source matrix S , with n by m elements, with target matrix T , with m by n elements.

Output Matrix T represents matrix S rotated 90°anti-clockwise.

```
procedure ROTATEMATRIXANTICLOCKWISE( $S, T, n, m$ )
     $f \leftarrow n - 1$ 
    for  $x = 0$  to  $n - 1$  do
        for  $y = 0$  to  $m - 1$  do
             $T[(f - x)][y] \leftarrow S[y][x]$  ▷  $S[\text{row}][\text{column}]$ 
        end for
    end for
end procedure
```

5.3.4.1 Description of working Algorithms 6 and 7 depict both clockwise and anti-clock array rotation respectively. Simply, both perform substitution of x and y values for the target array. An additional variable is needed to hold an offset, so that the new array offsets may be calculated correctly¹⁶. Figure 5.2 shows the clockwise rotation operation in effect.

¹⁶While this is not necessary, it make for slightly cleaner code.




Figure 5.2: Array Rotation

5.3.4.2 Implementation Listing 17 and 18 show implementations of clockwise and anticlockwise rotation respectively.

```
public void ArrayRotateClockWise(E[][] source, E[][] target, int width, int height){
    int factor = height-1;
    for (int y = 0; y < height; y++) {
        for (int x = 0; x < width; x++) {
            target[x][(factor-y)] = source[y][x];
        }
    }
}
```

Listing 17: Array Clockwise Rotation (Java)

```
public void ArrayAntiRotateClockWise(E[][] source, E[][] T, int width, int height){
    int factor = width-1;
    for (int x = 0; x < width; x++) {
        for (int y = 0; y < height; y++) {
            target[(factor-x)][y] = source[y][x];
        }
    }
}
```

Listing 18: Array Anti-Clockwise Rotation (Java)

Performance of both of these implementations will suffer due to access not being performed in a row by row fashion, which limits the viability of the above algorithms. However they should still be considered and used, when the many operations are performed on the resulting array, rather than single use post rotation.

5.3.4.3 Sample Problem- Matrix Rotation Given a two dimension matrix measuring $n \times m$, rotate the matrix in a clockwise direction.

INPUT

Input consists of one or more lines, with the first line being the size of the matrix to rotate in width and height. The width and height will be in a range between 0 and 79. The application should exit when the size given is "0 0".

The following lines denote the matrix to be rotated, with each element being a 1 or 0, separated by a space.

SAMPLE INPUT

```
3 3
1 0 1
0 0 0
1 0 0
2 2
1 0
0 1
0 0
```

OUTPUT

The output of each rotation should be the resultant matrix, without spaces between each column in the matrix. There must a blank line separating each resultant matrix.

SAMPLE OUTPUT

```
101
000
001

00
01
```

5.3.4.4 Sample Solution A solution to the sample problem utilising an array rotation can be seen in Listing 19. The solution performs the following:

1. Reads in the width and height values. If these are not 0 and 0, continue the main loop body, otherwise exit.
2. If either of the values is 0, (indicating a 0 width or height), skip attempting to read in a matrix, and go back to step 1.
3. Create a new array of size width, height to hold integers. Read in the values from stdin, and fill in these values into the respective position within the array.
4. Create a new array to hold the result of the rotation. Rotate the array, and print the results. Goto step 1.

As this challenge is very simple, the only issue to account for is if either the width or height is given as 0 (zero), in which case a blank line should be returned.

```
import java.util.Scanner;
public class ArrayRotation {

    /**
     * Rotate an Array in a clockwise direction
     *
     * @param source The source array
     * @param target The target array
     * @param width The width of the source array
     * @param height The height of the target array
     */
    public static void ArrayRotateClockWise(int[][] source, int[][] target,
                                             int width, int height) {
        int factor = height - 1;
        for (int y = 0; y < height; y++) {
            for (int x = 0; x < width; x++) {
                target[x][(factor - y)] = source[y][x];
            }
        }
    }
}
```

```
/**
 * Main
 */
public static void main(String[] args) {
    Scanner in = new Scanner(System.in);

    // get the size of the matrix to rotate.
    String matrixSizeLine = in.nextLine();
    Scanner matrixSize = new Scanner(matrixSizeLine);
    int width = matrixSize.nextInt();
    int height = matrixSize.nextInt();

    // Keep reading in a matrix until exit condition
    while (!((width == 0) && (height == 0))) {

        // Ignore the line if either width or height = 0.
        if (!((width == 0) || (height == 0))) {

            // Read in our matrix.
            int[][] matrix = new int[height][width];
            for (int count = 0; count < height; count++) {
                String[] matrixLine = in.nextLine().split("\\s+");
                for (int element = 0; element < matrixLine.length; element++) {
                    matrix[count][element] = Integer.parseInt(matrixLine[element]);
                }
            }

            // Rotate the matrix and print.
            int[][] target = new int[width][height];
            ArrayRotateClockWise(matrix, target, width, height);

            // Print the resultant matrix.
            for (int row = 0; row < width; row++) {
                for (int column = 0; column < height; column++) {
                    System.out.print(target[row][column]);
                }
                System.out.println();
            }

            // Print line between each matrix output
            System.out.println();
        }
        // get the size of the matrix to rotate.
        matrixSizeLine = in.nextLine();
        matrixSize = new Scanner(matrixSizeLine);
        width = matrixSize.nextInt();
        height = matrixSize.nextInt();
    }
}
```

Listing 19: Solution to Array Rotation Problem (Java)

5.3.5 Array Mirroring or Flipping

Array mirroring or flipping is a very simple technique, that simply requires a offset to be calculated from the current height or width value. Typically, to calculate the mirror location, we simple subtract the current source column (or row) value from the width (or height) size to calculate the target location. Algorithms 8 and 9 depict mirroring along the vertical axis and horizontal axis respectively.

Algorithm 8 Mirror Array Along Vertical Axis

Input A source matrix S , with n by m elements, with target matrix T , with n by m elements.**Output** Matrix T represents matrix S mirrored along a vertical axis

```
procedure MIRRORMATRIXVERTICAL( $S, T, n, m$ )  
  for  $y = 0$  to  $m - 1$  do  
    for  $x = 0$  to  $n - 1$  do  
       $T[m - y][x] \leftarrow S[y][x]$  ▷  $S[\text{row}][\text{column}]$   
    end for  
  end for  
end procedure
```

Algorithm 9 Mirror Array Along Horizontal Axis

Input A source matrix S , with n by m elements, with target matrix T , with n by m elements.**Output** Matrix T represents matrix S mirrored along a horizontal axis

```
procedure MIRRORMATRIXHORIZONTAL( $S, T, n, m$ )  
  for  $y = 0$  to  $m - 1$  do  
    for  $x = 0$  to  $n - 1$  do  
       $T[y][n - x] \leftarrow S[y][x]$  ▷  $S[\text{row}][\text{column}]$   
    end for  
  end for  
end procedure
```

As both of these algorithms access both the source and target arrays row by row, there should be no (or few) performance penalties due to cache misses.

6 Advanced Algorithms

6.1 Simple Maths

6.1.1 Greatest common divisor

The Greatest Common Divisor or Euclidean algorithm was originally developed by Euclid of Alexandria in 3rd BC, and computes the greatest common divisor of two non-negative, not-both-zero integers m and n , that is the largest integer that divides both m and n *evenly*.

6.1.1.1 Description of working The Euclidean algorithm is based on the principle that the greatest common divisor of two numbers does not change if the smaller number is subtracted from the larger number. For if k , m , and n are integers, and k is a common factor of two integers A and B , then $A = (n \times k)$ and $B = (m \times k)$ implies $A - B = (n - m) \times k$, therefore k is also a common factor of the difference. That k may also represent the greatest common divisor is proven below. For example, 21 is the GCD of 252 and 105 ($252 = 12 \times 21$; $105 = 5 \times 21$); since $252 - 105 = (12 - 5) \times 21 = 147$, the GCD of 147 and 105 is also 21.

Since the larger of the two numbers is reduced, repeating this process gives successively smaller numbers until one of them is zero. When that occurs, the GCD is the remaining nonzero number. By reversing the steps in the Euclidean algorithm, the GCD can be expressed as a sum of the two original numbers each multiplied by a positive or negative integer, e.g., $21 = [5 \times 105] + [(-2) \times 252]$.

Iterative and recursive implementations of the algorithm are shown in Algorithms 10 and 11 respectively.

Algorithm 10 Euclidean Algorithm (Iterative)

Input Positives integers m and n which may share a common divisor

Output Greatest Common Divisor, being $n \geq 1$

```
procedure EUCLID( $m, n$ )  
  while  $n \neq 0$  do                                     ▷ We have the answer if  $n$  is 0  
     $r \leftarrow m \bmod n$   
     $m \leftarrow n$   
     $n \leftarrow r$   
  end while  
  return  $m$                                              ▷ The gcd is  $m$   
end procedure
```

Algorithm 11 Euclidean Algorithm (Recursive)

```
procedure EUCLID( $m, n$ )  
  if  $n = 0$  then  
    return  $m$   
  end if  
  return gcd( $n, m \bmod n$ )  
end procedure
```

6.1.1.2 Implementation An implementation of the iterative form of the algorithm can be found in Listing 20.

```
public static int gcd(int m, int n) {
    int r;
    while (n != 0) {
        r = m % n;
        m = n;
        n = r;
    }
    return m;
}
```

Listing 20: Euclidean Algorithm (Java)

6.1.1.3 Sample Problem - GCD Calculate the GCD of two integers.**INPUT**

The first line of input will be N ($1 \leq N \leq 1000$), the number of test cases to run. On each of the next N lines will be two integers m and n ($0 \leq m, n \leq 2^{31} - 1$).

SAMPLE INPUT

```
3
12 60
60 24
3 5
```

OUTPUT

Output the greatest common denominator of m and n , one value on each line.

SAMPLE OUTPUT

```
12
12
1
```

6.1.1.4 Sample Solution The solution to this problem utilising Euclid's GCD algorithm can be seen in Listing 21.

Whilst this solution/challenge is rather simple, the following test cases must be considered:

1. One or both values are 0 (zero).
2. There will be longer run times if one or both numbers are primes. (The GCD is 1 in this case).
3. The input range is $0 < m < 2^{31} - 1$, which requires the use of at least a signed 32bit number (typically int on all platforms).

```
import java.util.Scanner;
public class GCD {
    /**
     * Compute the Greatest Common Divisor
     *
     * @param m number 1, greater or equal to 0.
     * @param n number 2, greater or equal to 0.
     * @return the GCD
     */
    public static int EuclidGCD(int m, int n) {
        int r;
        while (n != 0) {
            r = m % n;
```

```
        m = n;
        n = r;
    }
    return m;
}

/**
 * Main
 */
public static void main(String[] args) {
    Scanner in = new Scanner(System.in);

    // get first line and get the number of cases to test.
    int caseCount = Integer.parseInt(in.nextLine());

    // Keep reading each line while caseCount > 0
    while (caseCount-- > 0) {
        String line = in.nextLine();

        // extract numbers ints
        Scanner sc = new Scanner(line);
        int a = sc.nextInt();
        int b = sc.nextInt();

        // Output the gcd for the two given numbers.
        System.out.printf("%d\n", EuclidGCD(a,b));
    }
}
```

Listing 21: Solution to GCD Problem (Java)

6.1.2 Sieve of Eratosthenes (prime number generation)

The brute force method of finding prime numbers to limit N , can be found by checking for divisibility of all numbers between 2 and $N-1$. There are however more efficient methods for finding prime numbers, including the use of Sieve's, in particular the Sieve of Eratosthenes.

The Sieve of Eratosthenes, is a relatively simple algorithm, that performs prime number generation through iteration of composites based on the last prime found. This is done by maintaining an array of all numbers between 2 and N , and uses flags in each array slot to indicate if it's a prime, composite or unknown.

The Sieve of Eratosthenes is considered a $O(n \log \log n)$ operation, with a $O(n)$ memory requirement.

6.1.2.1 Description of working The algorithm utilises two arrays of numbers, one array that holds all numbers from 2 to N , and a second array that holds just the prime numbers from 2 to N .

The algorithm counts from 2 to $\lfloor \sqrt{N} \rfloor$, working on the array of all numbers from 2 to N . If the current field of the array is not 0 (that is, it's a known composite), it marks off all composites of the current number forward of the current position in the array of all numbers from 2 to N . Once it has cycled through the array, all the numbers that are left are primes. Figure 6.1 demonstrates this on a list of numbers from 1 to 30.

Algorithm 12 Sieve of Eratosthenes

Input Positives integer n denoting upper limit for prime number search, and vector S for list of prime numbers.**Output** All prime number $\leq n$ stored in vector S .

```
procedure ERATOSTHENES( $n, S$ )
  define vector  $A$  of size  $n$ 
  for  $p \leftarrow 2$  to  $n$  do
     $A[p] \leftarrow p$ 
  end for
  for  $p \leftarrow 2$  to  $\lfloor \sqrt{n} \rfloor$  do
    if  $A[p] \neq 0$  then
       $j \leftarrow p^2$ 
      while  $j \leq n$  do
         $A[j] \leftarrow 0$ 
         $j \leftarrow j + p$ 
      end while
    end if
  end for

   $i \leftarrow 0$  ▷ Copy list of primes from  $A$  to  $S$ 
  for  $p \leftarrow 2$  to  $n$  do
    if  $A[p] \neq 0$  then
       $S[i] \leftarrow A[p]$ 
       $i \leftarrow i + 1$ 
    end if
  end for
end procedure
```

First generate a list of integers from 2 to 30:

2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30
--

First number in the list is 2; cross out every 2nd number in the list after it (by counting up in increments of 2), i.e. all the multiples of 2:

2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30
--

Next number in the list after 2 is 3; cross out every 3rd number in the list after it (by counting up in increments of 3), i.e. all the multiples of 3:

2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30
--

Next number not yet crossed out in the list after 3 is 5; cross out every 5th number in the list after it (by counting up in increments of 5), i.e. all the multiples of 5:

2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30
--

Next number not yet crossed out in the list after 5 is 7; the next step would be to cross out every 7th number in the list after it, but they are all already crossed out at this point, as these numbers (14, 21, 28) are also multiples of smaller primes because 7×7 is greater than 30. The numbers left not crossed out in the list at this point are all the prime numbers below 30:

2 3 4 5 7 11 13 17 19 23 29

Figure 6.1: Demonstration of Sieve of Eratosthenes

6.1.2.2 Implementation An implementation of the Sieve of Eratosthenes can be found in Listing 22.

```
public static int[] Eratosthenes(int maxPrimeNumber){  
  
    // Define new vector A with all numbers from 2 to N  
    int[] vectorA = new int[maxPrimeNumber+1];  
    for(int index = 2; index <= maxPrimeNumber; index++){  
        vectorA[index] = index;  
    }  
  
    int nonPrimeCount = 0;  
    // Flag all non-primes in vector A.  
    int sqrtMaxPrime = (int)Math.sqrt(maxPrimeNumber);  
    for(int index = 2; index < sqrtMaxPrime; index++){  
        // Skip this number if known composite  
        if(vectorA[index] != 0){  
            int composite = index * index;  
            // Flag all composites of this number as non-prime  
            while(composite <= maxPrimeNumber){  
                if(vectorA[composite] != 0){  
                    vectorA[composite] = 0;  
                    nonPrimeCount++;  
                }  
                composite = composite + index;  
            }  
        }  
    }  
  
    // Create new array from A, filled with just primes.  
    int resultIndex = 0;  
    int[] primes = new int[maxPrimeNumber-nonPrimeCount-1];  
    for(int index = 2; index <= maxPrimeNumber; index++){  
        if(vectorA[index] != 0){  
            primes[resultIndex++] = vectorA[index];  
        }  
    }  
    return primes;  
}
```

Listing 22: Sieve of Eratosthenes Algorithm (Java)

6.1.2.3 Sample Problem - Primes for Hashtable Student Programmer Joe is developing a dynamic hash table class in Java, and from his textbooks knows that prime numbers make the best values to use within the hashing function. He would like a program to tell him if the number he is using is a prime, and if not what's the closest prime to but below to the number he has given the program.

INPUT

The input will be an integer indicating the number Joe inputs to be tested for being a prime, which is within the range of $0 \leq N \leq 2^{16} - 1$. A zero indicates the end of input. You should not process this input.

SAMPLE INPUT

```
3  
6243  
8191  
10  
0
```

OUTPUT

Output "prime" if a prime number, otherwise the closest prime below the number given.

SAMPLE OUTPUT

```
prime
6421
prime
7
```

6.1.2.4 Sample Solution A sample solution for the Primes of Hash Table problem can be found in Listing 23. There are 6542 prime numbers in the specified number space.

There are a number of methods to solve this problem, including:

1. Testing each number input as a prime using the divisibility test, and if not a prime decrementing the number by 1 and retesting until a prime is found.
2. Create a list of all numbers within the test space in an array and set all array elements that represent composites to 0. To test each number given, we simply look up the slot for that number in the array, and determine if it's a prime or not¹⁷. If it's not a prime, we use a reverse linear search through the array until a prime is found, and return the found prime. This offers, $O(1)$ performance for testing if a number is a prime within the array, but up to $O(71)$ if the number supplied is not a prime. The $O(71)$ is determined by the largest gap between prime numbers in the number space (being 71), and would be the worst case for a linear search operation if needed.
3. Create a list of all prime numbers in the test space, and store only these in an array¹⁷. Utilising a binary search to search for the number in the list, we can determine if the number is prime based on $O(\log_2 n)$ performance. With a slight modification to the binary search method, instead of returning -1 if not found, we can return the index of the last tested value, which incidentally will be the next prime below the number being tested. Therefore we can assume $O(\log_2 n)$ performance in the worst case for each number test. (With 6542 primes, this would result in $O(13)$ in the worst case).

For this problem, I utilised the third option listed above for three main reasons:

1. The memory space requirement is reduced after the lists of primes has been calculated¹⁸, allowing more of the list of numbers to fit into CPU cache. (Option 2 will utilise 256KB of RAM vs Option 3 with 26KB of RAM). While 256KB vs 26KB may not seem a lot, 26KB will fit or mostly fit into the Level 1 cache on most CPUs thereby increasing performance. If we expand the address space to the maximum value being $2^{31}-1$, Option 2 will utilise 8GB of RAM, whilst Option 3 will utilise 420MB of RAM once the lists have been created.
2. The time needed for determining if a prime is excellent with Option 2, but suffers some performance penalty if a linear search is needed to find the next lower prime. The largest gap in the required test space is 71 numbers, which may negatively effect performance. If we expand the address space to the maximum value being $2^{31}-1$, in Option 2 the largest gap between primes becomes 354 numbers, therefore requiring up to 354 operations for the linear search. Option 3 suffers little performance penalty, as the search remains at $O(\log_2 n)$, where $n = 105,097,565$, or $O(26)$ to determine if a prime or the next lower prime.
3. Utilising the Sieve of Eratosthenes, the resulting list of primes is naturally sorted allowing the use of a binary search to find the prime number, without additional work.

Option 1 was not considered due to the lack of CPU efficiency within the proposed solution.

The implemented solution performs the following steps:

1. Create a list of primes in an array utilising the Sieve of Eratosthenes algorithm. (If a complete array can not fit into memory, an alternate Segmented Sieve of Eratosthenes may be used to build the list of primes).

¹⁷This allows the search for prime to become a lookup operation rather than a calculation operation improving performance.

¹⁸Whilst during the creation of the list of primes, large amounts of memory is needed, once the list has been created, the needed space to hold the results is vastly smaller.

2. Reads in the next number to be tested. (If '0', exit).
3. Utilise a binary search to find the index of the number in the list of primes, and test the value stored against the number to determine if prime or not. If it is not a prime, the index points to the next prime lower than the number entered.

To test the implementation, a range of number were randomly selected to test the basic operation. To ensure that the implementation was robust, the following numbers were also tested:

- Number 1. Number 1 is neither a prime, nor composite. In this case, I return 1 due to 1 being non-prime.
- Number 2.
- Less than 0, which returns 2. (Less than 0, is invalid input, so can't be considering incorrect).
- More than 2^{16} , which the first prime less than 2^{16} .

Since a binary search was used, I had to pay little attention to values outside the input space, as the results of the binary search would always return an index that was within the address space. The exception to this was the number 1, in which a special case is used.

```
import java.util.Scanner;
public class PrimesForHashtable {
    /* Largest value with input space */
    final static int PRIME_SPACE = (int) Math.pow(2, 16);

    /**
     * Utilise Sieve of Eratosthenes to create an array of prime numbers
     */
    public static int[] Eratosthenes(int maxPrimeNumber) {
        // Define new vector A with all numbers from 2 to N
        int[] vectorA = new int[maxPrimeNumber + 1];
        for (int index = 2; index <= maxPrimeNumber; index++) {
            vectorA[index] = index;
        }
        int nonPrimeCount = 0;
        // Flag all non-primes in vector A.
        int sqrtMaxPrime = (int) Math.sqrt(maxPrimeNumber);
        for (int index = 2; index < sqrtMaxPrime; index++) {
            // Skip this number if known composite
            if (vectorA[index] != 0) {
                int composite = index * index;
                // Flag all composites of this number as non-prime
                while (composite <= maxPrimeNumber) {
                    if (vectorA[composite] != 0) {
                        vectorA[composite] = 0;
                        nonPrimeCount++;
                    }
                    composite = composite + index;
                }
            }
        }
        // Create new array from A, filled with just primes.
        int resultIndex = 0;
        int[] primes = new int[maxPrimeNumber - nonPrimeCount - 1];
        for (int index = 2; index <= maxPrimeNumber; index++) {
            if (vectorA[index] != 0) {
                primes[resultIndex++] = vectorA[index];
            }
        }
        return primes;
    }
}
```

```
}

/**
 * Perform binary search of array to locate prime number.
 */
public static int BinarySearch(int[] vector, int key, int low, int high) {
    if (low > high) {
        return (high >= 0 ? high : 0);
        // Return the index of the current high value, as this
        // points to 1 below the key, if not found
    }
    int mid = (low + high) / 2;
    if (vector[mid] == key) {
        return mid;
    } else {
        if (vector[mid] > key) {
            return BinarySearch(vector, key, low, (mid - 1));
        } else {
            return BinarySearch(vector, key, (mid + 1), high);
        }
    }
}

/**
 * Returns a prime equal to or below to the number given.
 */
public static int LookUpPrime(int[] primes, int number) {
    return primes[BinarySearch(primes, number, 0, primes.length-1)];
}

/**
 * Main
 */
public static void main(String[] args) {

    // Create a list of primes for the given input space.
    int[] primes = Eratosthenes(PRIME_SPACE);

    Scanner in = new Scanner(System.in);
    String line = in.nextLine();

    // Keep processing input until 0.
    while (!line.equals("0")) {
        // Get our number, and it's closest prime.
        int number = Integer.parseInt(line);
        // Handle special case of number 1.
        if (number != 1) {
            int prime = LookUpPrime(primes, number);
            // Display if it's a prime, or closet prime.
            if (prime == number) {
                System.out.println("prime");
            } else {
                System.out.println(prime);
            }
        } else {
            System.out.println("1");
        }
        line = in.nextLine();
    }
}
```

Listing 23: Solution to Primes for Hashtable (Java)

6.2 String based algorithms and data structures

String are a fundamental data type with a high importance in the following areas:

- Web crawlers utilised by search engines classify billions of text documents each day.
- Bioinformatics Research, particularly around genome based research.

Strings are effectively a collection of characters, with each character being defined as part of an alphabet. This alphabet, may consist of only a single character, a small number of characters, (eg ASCII with 127 characters), a large number of characters (eg Unicode with 1,114,112 characters), or all possible values represented by either a byte (256 characters), word (65536 characters) or double-word (2^{32} characters). One common misconception is that a string may only represent those letters and figures as defined in a human language, but from a Computers viewpoint any value that can be defined may be included (or even excluded) from the alphabet that defines a string.

Typically, we allow A to be a alphabet (finite set), with:

- *pattern* and *text* are vectors of elements of A .
- A may be:
 - usual human alphabet
 - binary alphabet = $A = \{0,1\}$
 - DNA alphabet = $A = \{A, C, G, T\}$
 - etc.

When dealing with strings in different languages, the meaning and implementation may vary.

- C99 - A string is defined as a vector of char (or byte), utilising ASCII¹⁹ or UTF-8²⁰ encoding terminated by the character *NUL* (0).
- C++ - A string may be that as C99, or of an object wrapped around a C99 string. Some operating systems like Windows allow a Wide-Character (16bits) to be utilised in place of single byte characters.
- Java - A string is a vector of UTF-16 encoded characters, represented as an object, and is immutable.
- C# (with .NET) - A string is a vector of UTF-16 encoded characters, represented as an object, and is immutable.

There are numerous algorithms that perform string matching, including Brute Force, Rabin-Karp, Knuth-Morris-Pratt, Boyer-Moore, and Horspool. Regular Expressions are also considered a form of pattern matching, with the advantage that they allow an infinite number of patterns to be matched.

This guide will discuss the Brute Force and Knuth-Morris-Pratt algorithms.

6.2.1 Brute Force Substring Search

Brute Force Substring search is considered a naive based algorithm that utilises a very simple method to determine string or pattern matching.

The algorithm walks through the vector of text, attempting to match the pattern in sequence.

¹⁹Other non-PC systems may utilise alternate encoding systems such as EBCDIC in place of ASCII encoding.

²⁰Modern UNIX implementations like FreeBSD, Solaris and GNU/Linux use UTF-8.

6.2.1.2 Implementation An implementation of the Brute Force String Matching can be found in Listing 24.

```
int BruteForceStringMatch(char* text, int textLength,
                          char* pattern, int patternLength){
    int patternIndex = 0;
    // Scan the text looking for the pattern
    for(int textIndex = 0; textIndex <= textLength-patternLength; textIndex++){
        patternIndex = 0; // Reset index into pattern
        // While the current text position matches the pattern, keep scanning it.
        while((patternIndex < patternLength)
              &&(pattern[patternIndex] == text[textIndex+patternIndex])){
            patternIndex++;
        }
        // If the patternIndex is equal to patternLength, we have found the pattern!
        if(patternIndex == patternLength){
            return textIndex;
        }
    }
    // Pattern not found, so return -1.
    return -1;
}
```

Listing 24: Brute Force String Matching (C++)

6.2.1.3 Sample Problem - Where's Waldorf Given a m by n grid of letters, ($1 \leq m, n \leq 50$), and a list of words, find the location in the grid at which the word can be found. A word matches a straight, uninterrupted line of letters in the grid. A word can match the letters in the grid regardless of case (i.e. upper and lower case letters are to be treated as the same). The matching can be done in any of the eight directions either horizontally, vertically or diagonally through the grid.

INPUT

The input begins with a single positive integer on a line by itself indicating the number of the cases following, each of them as described below. This line is followed by a blank line, and there is also a blank line between two consecutive inputs.

The input begins with a pair of integers, m followed by n , $1 \leq m, n \leq 50$ in decimal notation on a single line. The next m lines contain n letters each; this is the grid of letters in which the words of the list must be found. The letters in the grid may be in upper or lower case.

Following the grid of letters, another integer k appears on a line by itself ($1 \leq k \leq 20$). The next k lines of input contain the list of words to search for, one word per line. These words may contain upper and lower case letters only (no spaces, hyphens or other non-alphabetic characters).

SAMPLE INPUT

```
1

8 11
abcDEFGhigg
hEbkWalDork
FtyAwaldORm
FtsimrLqsrc
byoArBeDeyv
Klcbqwikomk
strEBGadhrb
yUiqlxcnBjf
4
Waldorf
Bambi
Betty
Dagbert
```


OUTPUT

For each test case, the output must follow the description below. The outputs of two consecutive cases will be separated by a blank line.

For each word in the word list, a pair of integers representing the location of the corresponding word in the grid must be output. The integers must be separated by a single space. The first integer is the line in the grid where the first letter of the given word can be found (1 represents the topmost line in the grid, and m represents the bottommost line). The second integer is the column in the grid where the first letter of the given word can be found (1 represents the leftmost column in the grid, and n represents the rightmost column in the grid). If a word can be found more than once in the grid, then the location which is output should correspond to the uppermost occurrence of the word (i.e. the occurrence which places the first letter of the word closest to the top of the grid). If two or more words are uppermost, the output should correspond to the leftmost of these occurrences. All words can be found at least once in the grid.

SAMPLE OUTPUT

```
2 5
2 3
1 2
7 8
```

6.2.1.4 Sample Solution A sample solution to the Where's Waldorf problem can be found in Listing 25.

```
#include <cstdlib>
#include <iostream>
#include <cctype>
#include <cstring>

using namespace std;

char** wordmap;
int testCases = 0;
int wordmapWidth = 0;
int wordmapHeight = 0;
char blankline[256];
int wordSearches = 0;
char* pattern = NULL;
int patternLength = 0;
char* text = NULL;
int location = -1;
int currentRow = 0;
int currentCol = 0;

/**
 * Attempt to find a pattern contained within the text.
 */
int BruteForceStringMatch(char* text, int textLength, char* pattern, int patternLength) {
    int patternIndex = 0;

    // Ensure the text is big enough to hold our pattern!
    if (patternLength > textLength) {
        return -1;
    }

    // Scan the text looking for the pattern
    for (int textIndex = 0; textIndex <= textLength - patternLength; textIndex++) {
        patternIndex = 0;
        // While the current text position matches the pattern, keep scanning it.
        while ((patternIndex < patternLength) && (pattern[patternIndex] ==
            text[textIndex + patternIndex])) {
            patternIndex++;
        }
        // If the patternIndex is equal to the patternLength, we have found the pattern!
        if (patternIndex == patternLength) {
            return textIndex;
        }
    }
}
```

```
}
// Pattern not found, so return -1.
return -1;
}

/**
 * Create a search pattern of length, from the word map.
 * @param col starting position in the word map (0 offset)
 * @param row starting position in the word map (0 offset)
 * @param length length of required string.
 * @param colOffset direction to move column index for each letter.
 * @param rowOffset direction to move row index for each letter.
 * @return Index of found string.
 */
char* buildTextString(int col, int row, int length, int colOffset, int rowOffset) {
    char* newString = new char[length + 1];
    int index = 0;
    while (length--) {
        newString[index++] = wordmap[row][col];
        row += rowOffset;
        col += colOffset;
        // Ensure we are not going outside map boundary
        if ((row < 0) || (row >= wordmapHeight) || (col < 0) || (col >= wordmapWidth)) {
            newString[index] = NULL;
            return newString;
        }
    }
    newString[index] = NULL;
    return newString;
}

/**
 * Convert string to lower case, and store in place
 * @param str ASCII string to convert.
 */
void strToLower(char* str) {
    int i = 0;
    while (str[i]) {
        str[i] = (char) tolower(str[i]);
        i++;
    }
}

/**
 * Update the best found location, based on current row/col positions
 */
void UpdateFind(int row, int col, int rowOffset, int colOffset, int location) {
    if (currentRow > row + (rowOffset * location)) {
        currentRow = row + (rowOffset * location);
        currentCol = col + (colOffset * location);
    } else if (currentRow == row + (rowOffset * location)) {
        if (currentCol > col + (colOffset * location)) {
            currentRow = row + (rowOffset * location);
            currentCol = col + (colOffset * location);
        }
    }
}

/**
 * Main
 */
int main() {
    scanf("%d", &testCases);
    while (testCases--) {
        gets(blankline);
        // Read in the wordmap, and ensure in lower case.
        scanf("%d_%d", &wordmapHeight, &wordmapWidth);
        wordmap = new char[wordmapHeight];
        for (int row = 0; row < wordmapHeight; row++) {
            wordmap[row] = new char[wordmapWidth + 1];
            scanf("%s", wordmap[row]);
            strToLower(wordmap[row]);
        }
    }
}
```

```
// Read in the number of words to search for.
scanf("%d", &wordSearches);
while (wordSearches--) {
    // Read in a single word.
    pattern = new char[wordmapWidth + 1];
    scanf("%s", pattern);
    strToLower(pattern);
    patternLength = strlen(pattern);

    currentRow = wordmapHeight;
    currentCol = wordmapWidth;

    //Attempt to scan for word, along the top of the grid.
    for (int col = 0; col < wordmapWidth; col++) {
        int row = 0;
        int rowOffset = 1;
        for (int colOffset = -1; colOffset < 2; colOffset++) {
            // Get our next text string, and attempt to match!
            text = buildTextString(col, row, 50, colOffset, rowOffset);
            location = BruteForceStringMatch(text, strlen(text), pattern, patternLength);
            if (location >= 0) {
                UpdateFind(row, col, rowOffset, colOffset, location);
            }
            delete(text);
        }
    }

    //Attempt to scan for word, along the bottom of the grid.
    for (int col = 0; col < wordmapWidth; col++) {
        int row = wordmapHeight - 1;
        int rowOffset = -1;
        for (int colOffset = -1; colOffset < 2; colOffset++) {
            // Get our next text string, and attempt to match!
            text = buildTextString(col, row, 50, colOffset, rowOffset);
            location = BruteForceStringMatch(text, strlen(text), pattern, patternLength);
            if (location >= 0) {
                UpdateFind(row, col, rowOffset, colOffset, location);
            }
            delete(text);
        }
    }

    //Attempt to scan for word, along the left of the grid.
    for (int row = 0; row < wordmapHeight; row++) {
        int col = 0;
        int colOffset = 1;
        for (int rowOffset = -1; rowOffset < 2; rowOffset++) {
            // Get our next text string, and attempt to match!
            text = buildTextString(col, row, 50, colOffset, rowOffset);
            location = BruteForceStringMatch(text, strlen(text), pattern, patternLength);
            if (location >= 0) {
                UpdateFind(row, col, rowOffset, colOffset, location);
            }
            delete(text);
        }
    }

    //Attempt to scan for word, along the right of the grid.
    for (int row = 0; row < wordmapHeight; row++) {
        int col = wordmapWidth - 1;
        int colOffset = -1;
        for (int rowOffset = -1; rowOffset < 2; rowOffset++) {
            // Get our next text string, and attempt to match!
            text = buildTextString(col, row, 50, colOffset, rowOffset);
            location = BruteForceStringMatch(text, strlen(text), pattern, patternLength);
            if (location >= 0) {
                UpdateFind(row, col, rowOffset, colOffset, location);
            }
            delete(text);
        }
    }
}
```

```

    printf("%d_%d\n", currentRow + 1, currentCol + 1);
    // Free our word.
    delete(pattern);
}

// Free our wordmap;
for (int row = 0; row < wordmapHeight; row++) {
    delete(wordmap[row]);
}
delete(wordmap);
printf("\n");
}
return 0;
}

```

Listing 25: Solution to Where's Waldorf (C++)

The Where's Waldorf problem, can be solved a number of ways, with the two most common methods being:

1. Move through each letter in the grid, and attempt to find the word in one of the eight directions leading from the current character, until a match is found. (This is roughly $O(n^2)$ performance).
2. Move along the outer most characters only, and attempt to find the word in three directions that look inward into the grid, utilising a pattern search algorithm. (This roughly $O(12n)$ performance).

For solving this problem, I utilised the second method.

In order to mitigate the case from being an issue, when reading in either a grid line or pattern to look for, I convert both to lower case before storing/using them. The sample solution first initiates a scan along the top of the grid, creating 3 texts to search for each position along the top of the grid. Figure 6.3 demonstrates this technique. If found, I update the best known location (according to current row and column values), and then continue to scan the rest of the grid (as shown as the gradients squares in Figure 6.3). Lastly at the end of the grid scan, I output the best found location.




Figure 6.3: Scanning Grid for Where's Waldorf

Some of the tough test cases include:

- Grids being either 1 wide or 1 row deep.
- Current text to scan is smaller than the pattern length.
- Multiple words found in the grid.
- A grid being all of one letter, and pattern to find is a string of the one letter.

6.2.2 Knuth-Morris-Pratt Substring Search

When studying the worst-case performance of the brute force pattern matching, there is a major deficiency in that characters are potentially rescanned. Specifically, we may perform many comparisons while testing potential placement of the pattern against the text, yet if we discover a pattern character that does not match in the text, then we throw away all the information gained by these comparisons and start over again from scratch with the algorithm.

The Knuth-Morris-Pratt (or KMP) algorithm, avoids this waste of information, and in doing so, achieves $O(n+m)$ performance, which is optimal in the worst case.

Algorithm 14 Knuth-Morris-Pratt Substring Search

Input A Vector T of n characters representing a text, and vector P of m characters representing a pattern to find in T .

Output The index of the first character of P when a complete match is found in T , otherwise -1 if search was unsuccessful.

```
procedure KMPMATCH( $T, P$ )
   $f \leftarrow \text{KMFAILURE}(P)$ 
   $i \leftarrow 0$ 
   $j \leftarrow 0$ 
  while  $i < n$  do
    if  $P[j] = T[i]$  then
      if  $j = m - 1$  then
        return  $i - m + 1$                                 ▷ A match!
      end if
       $i \leftarrow i + 1$ 
       $j \leftarrow j + 1$ 
    else if  $j > 0$  then                                    ▷ No match, but advanced in P
       $j \leftarrow f(j - 1)$ 
    else
       $i \leftarrow i + 1$ 
    end if
  end while
  return -1
end procedure
```

Input A Vector P of m characters representing a pattern.

Output The failure function f for P , which maps j to the length of the longest prefix of P that is a suffix of $P[1..j]$

```
procedure KMFAILURE( $P$ )
   $i \leftarrow 1$ 
   $j \leftarrow 0$ 
   $f(0) \leftarrow 0$ 
  while  $i < m$  do
    if  $P[j] = P[i]$  then                                    ▷ We have matched  $j + 1$  characters
       $f(i) \leftarrow j + 1$ 
       $i \leftarrow i + 1$ 
       $j \leftarrow j + 1$ 
    else if  $j > 0$  then                                      ▷  $j$  indexes just after a prefix of  $P$  that must match
       $j \leftarrow f(j - 1)$ 
    else                                                    ▷ We have no matches
       $f(i) \leftarrow 0$ 
       $i \leftarrow i + 1$ 
    end if
  end while
end procedure
```

6.2.2.1 Description of working The main idea of the KMP algorithm is to preprocess the pattern string P so as to compute a **failure function** f that indicates the proper shift of P so that, to the largest extent possible, we can reuse previously performed comparisons. Specifically, the failure function $f(j)$ is defined as the length of the longest prefix of P that is a suffix of $P[1..j]$. We also use the convention that $f(0) = 0$. The importance of this failure function is that “encodes” repeated substrings inside the pattern itself.

The KMP pattern matching algorithm, shown in Algorithm 14, incrementally processes the text string T comparing it to the pattern string P . Each time there is a match, we increment the current indices. On the other hand, if there is a mismatch and we have previously made progress in P , then we consult the failure function to determine the new index in P where we need to continue checking P against T . Otherwise, we simply increment the index for T (and keep the index variable for P at its beginning). We repeat this process until we find a match of P in T or the index of T reaches n , the length of T .

The main part of the KMP algorithm is the while-loop, which performs a comparison between a character in T and a character in P each iteration. Depending upon the outcome of this comparison, the algorithm either moves on to the next characters in T and P , consults the failure function for a new candidate character in P , or starts over with the next index in T . The correctness of this algorithm follows from the definition of the failure function. The skipped comparisons are actually unnecessary, for the failure function guarantees that all the ignored comparisons are redundant - they would involve comparing characters we already know match.

In Figure 6.4, we illustrate the execution of the KMP algorithm. The failure function f for the pattern displayed is given in Figure 6.5. The algorithm performs 19 character comparisons, which are indicated with numerical labels. Note the use of the failure function to avoid redoing one of the comparisons between a character of the pattern and a character of text. Also note that this algorithm performs fewer overall comparisons than the brute force algorithm.

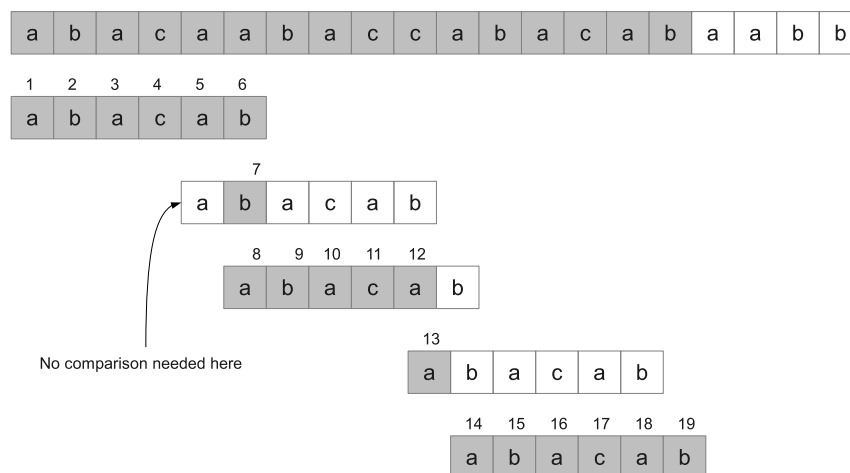


Figure 6.4: Illustration of the KMP algorithm.

Consider the pattern string $P = \text{"abacab"}$. The failure function $f(j)$ for the string is shown as the following table.

j	0	1	2	3	4	5
$P(j)$	a	b	a	c	a	b
$f(j)$	0	0	1	0	1	2

Figure 6.5: Failure Function for Figure 6.4

The failure function utilised within the KMP Algorithm is also shown in Algorithm 14. This algorithm is another example of a bootstrapping process quite similar to that used in the KMPMATCH procedure. We

compare the pattern to itself as in the KMP Algorithm. Each time we have two characters that match, we set $f(i) = j + 1$. Note that since we have $i > j$ throughout the execution of the algorithm, $f(j - 1)$ is always defined when we need to use it.

6.2.2.2 Implementation An implementation of the KMP Algorithm can be found in Listing 26.

```
/**
 * KMP Failure function.
 * @param Pattern Pointer to the pattern to be processed to derive the failure
 *          function.
 * @param patternLength Length of pattern.
 * @return The failure function (array of int).
 */
int *KMPSuccess(char *Pattern, int patternLength) {
    int indexI = 1;
    int indexJ = 0;
    int* failure = new int[patternLength]; // Create new function f.

    failure[0] = 0;
    while (indexI < patternLength) {
        if (Pattern[indexJ] == Pattern[indexI]) {
            failure[indexI++] = ++indexJ;
        } else if (indexJ > 0) {
            indexJ = failure[indexJ - 1];
        } else {
            failure[indexI++] = 0;
        }
    }
    return failure;
}

/**
 * Utilise the KMP Substring search algorithm to find a pattern in text.
 * @param Text Pointer to the text to be searched.
 * @param textLength Length of the text to be processed.
 * @param Pattern Pointer to the pattern to find within the text
 * @param patternLength Length of the pattern string
 * @return Index of pattern in text, otherwise -1 if not found.
 */
int KMPSuccess(char* Text, int textLength, char* Pattern, int patternLength) {
    char* failure = KMPSuccess(Pattern, patternLength);
    int indexI = 0;
    int indexJ = 0;
    while (indexI < textLength) {
        if (Pattern[indexJ] == Text[indexI]) {
            if (indexJ == (patternLength - 1)) {
                delete(failure); // Found a match, so return the index.
                return (indexI - patternLength + 1);
            }
            indexI++;
            indexJ++;
        } else if (indexJ > 0) { // No match but advanced in P.
            indexJ = failure[indexJ - 1]; // Move forward appropriate amount.
        } else {
            indexI++;
        }
    }
    delete(failure);
    return -1;
}
```

Listing 26: KMP String Match (C++)

6.2.2.3 Sample Problem - Big String Search Search some text to locate a particular pattern. If the pattern exists in the text, output the 0-based index of the pattern in the text.

INPUT

The first line of input will be N ($1 \leq N \leq 1000$), the number of test cases to run. On each of the next $2N$ lines will be two strings. The first string will be the pattern to search for. On the next line will be the text to search in. All pattern strings will have a length between 1 and 100,000 inclusive. All text strings will have a length between 1 and 500,000 inclusive.

SAMPLE INPUT

```
5
the
the quick brown fox jumps over the lazy dog
000111
010101010101000110000111
CGAT
TGATCTAGCTAGCTAGCTAGCTAGCATACGCATAGCTA
happy
I could be happy if I didn't have to work
```

OUTPUT

For each test case, output the test case number followed by a space, and then the 0-based index of the first letter of the pattern in the text if the pattern can be found, or NOT FOUND if the pattern does not exist in the text.

SAMPLE OUTPUT

```
1 0
2 18
3 NOT FOUND
4 11
```

6.2.2.4 Sample Solution A sample solution to the Big String Search problem can be found in Listing 27.

```
#include <cstdlib>
#include <cstdio>
#include <cstring>

using namespace std;

#define MAX_PATTERN_LENGTH 100000
#define MAX_TEXT_LENGTH 500000

int testCases = 0;
char pattern[MAX_PATTERN_LENGTH];
char text[MAX_TEXT_LENGTH];

int index = 0;
int testCount = 1;

/**
 * KMP Failure function.
 * @param Pattern Pointer to the pattern to be processed to derive the failure function.
 * @param patternLength Length of pattern.
 * @return The failure function (array of int).
 */
int *KMPFailure(char *Pattern, int patternLength) {
    int indexI = 1;
    int indexJ = 0;
    int* failure = new int[patternLength]; // Create new function f.

    failure[0] = 0;
    while (indexI < patternLength) {
        if (Pattern[indexJ] == Pattern[indexI]) {
            failure[indexI++] = ++indexJ;
        }
    }
}
```



```
        } else if (indexJ > 0) {
            indexJ = failure[indexJ - 1];
        } else {
            failure[indexI++] = 0;
        }
    }
    return failure;
}

/**
 * Utilise the KMP Substring search algorithm to find a pattern in text.
 * @param Text Pointer to the text to be searched.
 * @param textLength Length of the text to be processed.
 * @param Pattern Pointer to the pattern to find within the text
 * @param patternLength Length of the pattern string
 * @return Index of pattern in text, otherwise -1 if not found.
 */
int KMPMatch(char* Text, int textLength, char* Pattern, int patternLength) {
    int* failure = KMPFailure(Pattern, patternLength);
    int indexI = 0;
    int indexJ = 0;
    while (indexI < textLength) {
        if (Pattern[indexJ] == Text[indexI]) {
            if (indexJ == (patternLength - 1)) {
                delete [] failure;
                return (indexI - patternLength + 1);
            }
            indexI++;
            indexJ++;
        } else if (indexJ > 0) { // No match but advanced in P.
            indexJ = failure[indexJ - 1]; // Move forward appropriate amount.
        } else {
            indexI++;
        }
    }
    delete [] failure;
    return -1;
}

/*
 * Main Function
 */
int main() {
    scanf("%d\n", &testCases);
    while (testCases--) {
        // Read in our pattern and text
        gets(pattern);
        gets(text);

        // found our pattern in text and report the location.
        index = KMPMatch(text, strlen(text), pattern, strlen(pattern));
        if (index == -1) {
            printf("%d NOT FOUND\n", testCount++);
        } else {
            printf("%d %d\n", testCount++, index);
        }
    }
    return 0;
}
```

Listing 27: Solution to Big String Search (C++)

The Big String Problem is a relatively simple problem that relies on the string search algorithm to be correct. As can be seen in Listing 27, within the main() function, while-loop comprises of simply reading in the pattern and text strings, performing the search function and reporting the results of the search.

Sample test cases include:

- Pattern and Text length being 1.
- Pattern length being greater than the text length.

6.3 Graph Theory (Basic)

Graph Theory forms one of the principle areas of modern computer science due to its ability to be applied to a wide range of problems and areas of study.

When viewed abstractly, a **graph** G is simply a set of V of **vertices** and a collection of E of pairs of vertices from V , called **edges**. Thus, a graph is a way of representing connections or relationships between pairs of objects from some set V . Edges in a graph are either **directed** or **undirected**, that is an edge may only allow connection in a single direction between vertices. Edges may also have a **weight** applied that defines the cost²¹ of utilising the edge for some defined operation.

Using the graph in Figure 6.6 as an example, each location (town/suburb) is a vertex, and the road is an edge between locations. The directions denote if the road is a one-way street or two-way, and the weight is the defined time to travel on that road to get between the vertices.

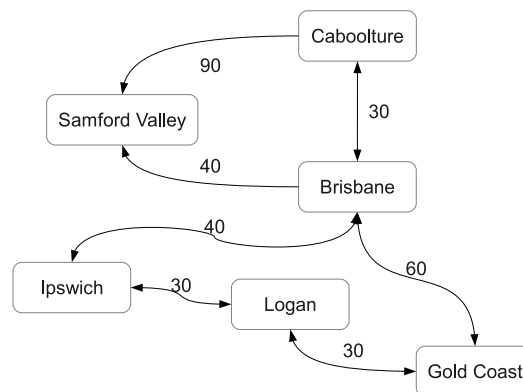


Figure 6.6: Basic Graph

Using the graph as shown, some examples include:

- It's quicker to travel from Caboolture to Samford Valley via Brisbane, than it is to go via a direct route. (70 vs 90)
- For someone living in Logan, in order to travel to Brisbane, must travel either through Ipswich or the Gold Coast. (via Ipswich offers the shortest time).
- The only way to travel from the Gold Coast to Samford Valley is via Brisbane.
- And if you arrive in Samford Valley, there is no road that leads out!

Other terminology that is often applied to Graph Theory includes:

- A **degree** of V , denoted $deg(v)$, is the number of edges incident to V . This may also be broken down into in-degree and ou-degree when applied with a directed graph. (The count of the edges leading to or from the vertex respectively).
- A **trail** in a graph, is a sequences of necessarily distinct vertices, but distinct edges.
- A **path** is a trail in which all vertices are distinct.
- A **circuit** is a trail in which the starting vertex is the same as the ending vertex.
- A **cycle** is a path in which the starting vertex is the same as the ending vertex.
- A **Hamiltonian cycle** is a cycle that contains all the vertices of a graph.
- A **sparse** graph is a graph with few edges, conversely, a **dense** graph is a graph with many edges.

Other terms commonly used within Graph Theory will be introduced later through out this guide as needed, especially when utilised within description of an algorithm.

²¹A weight may have either a positive or negative value depending on what is being modelled.

6.3.1 Representation as adjacency matrix and adjacency list.

Graphs may be represented in two forms, either as an adjacency matrix or as an adjacency list.




Figure 6.7: Basic Graph Example (weighted, undirected)

6.3.1.1 Adjacency List An Adjacency List is a method of representing a graph in a space efficient form, in that a vector (array) of vertices is maintained, and each element of the vector maintains a list (typically a linked list or ArrayList) of vertices that may be reached from the selected element. For weighted graphs, each element in the list also includes the edge weight in addition to the vertex.

Using the graph in Figure 6.7, the represented Adjacency List is shown in Table 3.

Vertex	List
A	B,5 → D,10
B	A,5 → C,6 → D,3 → E,2
C	B,6 → D,7
D	A,10 → B,3 → E,4 → F,8
E	B,2 → C,7 → D,4 → F,12 → G,6
F	D,8 → E,12 → G,1
G	E,6 → F,1

Table 3: Adjacency List

Two advantages of utilising the Adjacency List to represent a graph are:

- The space required to store the list is minimal in that no wasted space is present.
- Traversal of the graph is easy, by simply following the vertex information in the list.

6.3.1.1.1 Description of working The adjacency list works primarily as it stores references and not copies to other vertices within the list, as well as the weights associated with the edge to the vertex. Additionally the data structure naturally supports both unweighted and weighted graphs, as well as directed and undirected graphs within the one implementation.

For graph traversal, you can simply make one of vertices within the list the current location, and therefore find its edge list as well. This allows for a naturally recursive implementation for many graph traversal algorithms.

The disadvantages of the adjacency list implementation, is that in order to find if either an edge between vertices exist, or any cost associated, the list must be scanned to find the other vertex, which in worst case may be $O(n)$.

6.3.1.1.2 Implementation A storage class implemented in Java for representing an Adjacency List is shown in Listing 28.

```
import java.util.ArrayList;

public class AdjacencyList {

    /**
     * Define generic Vertex class to hold our information.
     */
    public class Vertex<V> {
        public V node;
        public Edge edgeList;
    }

    /**
     * Generic Edge List to store our edge information.
     */
    public class Edge{
        public Vertex node = null;
        public int weight = 0;
        public Edge next = null;
    }

    /**
     * The Adjacency list.
     */
    public Vertex<V>[] adjacencyList;
}
```

Listing 28: Adjacency List (Java)

6.3.1.1.3 Sample Problem - Adjacency List Construct the adjacency list representation of an unweighted undirected graph.

INPUT

The first line of input will be a number on a line by itself which is the number of test cases to run. For each test case, the first line will be two numbers separated by a space N and M , where N ($1 \leq N \leq 5000$) is the number of nodes in the graph and M ($1 \leq M \leq 10000$) is the number of edges. The graph nodes will be numbered 0 to $N-1$. Each of the next M lines contain 2 numbers A and B separated by a space representing an edge in the graph between A ($0 \leq A < N$) and B ($0 \leq B < N$).

SAMPLE INPUT

```
2
4 4
0 1
1 2
2 0
2 3
4 4
0 3
1 0
2 0
2 3
```

OUTPUT

For each test case, output the test case number on a line followed by the adjacency list representation of the graph. For each node n in the graph output a line listing the number of the node followed by a colon and a space, then a comma and single space separated list of the nodes adjacent to n . The adjacent nodes should be listed in ascending order. See the example output below.

SAMPLE OUTPUT

```
1
0: 1, 2
1: 0, 2
2: 0, 1, 3
3: 2
2
0: 1, 2, 3
1: 0
2: 0, 3
3: 0, 2
```

6.3.1.1.4 Sample Solution A sample solution to the Adjacency List problem can be found in Listing 29.

```
import java.util.ArrayList;
import java.util.Collections;
import java.util.Scanner;

public class AdjacencyList {

    /**
     * Define generic Vertex class to hold our information.
     */
    public static class Vertex {

        public int vertexID;
        public Edge edgeList;
        // Insert additional information related to the Vertex here.

        /**
         * Adds the edge into the adjacency list.
         *
         * @param source The source vertex ID
         * @param destination The destination vertex ID
         */
        public void addEdge(Vertex[] adjacencyList, int source, int destination) {

            // Get edge list for source.
            Edge edge = adjacencyList[source].edgeList;
            if (edge != null) {
                while (edge.next != null && edge.node.vertexID != destination) {
                    edge = edge.next;
                }
                if (edge.node.vertexID != destination) {
                    edge.next = new Edge(adjacencyList[destination], 0, edge.next);
                }
            } else {
                // Doesn't exist, so add it in.
                adjacencyList[source].edgeList =
                    new Edge(adjacencyList[destination], 0, null);
            }

            // Repeat same in reverse, as this is an undirected graph.
        }
    }
}
```

```
        edge = adjacencyList[destination].edgeList;
        if (edge != null) {
            while (edge.next != null && edge.node.vertexID != source) {
                edge = edge.next;
            }
            if (edge.node.vertexID != source) {
                edge.next = new Edge(adjacencyList[source], 0, edge.next);
            }
        } else {
            // Doesn't exist, so add it in.
            adjacencyList[destination].edgeList =
                new Edge(adjacencyList[source], 0, null);
        }
    }
}

/**
 * Generic Edge List to store our edge information.
 */
public static class Edge {

    public Vertex node = null;
    public int weight = 0;
    public Edge next = null;

    public Edge(Vertex vertex, int weight, Edge edge) {
        node = vertex;
        this.weight = weight;
        next = edge;
    }
}

/**
 * Main
 */
public static void main(String[] args) {

    Vertex[] adjacencyList;

    Scanner in = new Scanner(System.in);

    // get first line and get the number of cases to test.
    int caseCount = Integer.parseInt(in.nextLine());
    int loopCount = 0;
    while (caseCount - loopCount > 0) {
        String line = in.nextLine();

        // extract numbers, this is the node count and edge count.
        Scanner sc = new Scanner(line);
        int vertexCount = sc.nextInt();
        int edgeCount = sc.nextInt();

        // Create the vector of Vertices, and set the vertex number in each Vertex
        adjacencyList = new Vertex[vertexCount];
        for (int i = 0; i < vertexCount; i++) {
            adjacencyList[i] = new Vertex();
            adjacencyList[i].vertexID = i;
        }

        while (edgeCount-- > 0) {
            // Get our next edge...
            line = in.nextLine();

```

```
        // extract numbers, this is the node count and edge count.
        sc = new Scanner(line);
        int source = sc.nextInt();
        int destination = sc.nextInt();
        adjacencyList[0].addEdge(adjacencyList, source, destination);
    }

    System.out.printf("%d\n", ++loopCount);
    // Print the resulting adjacency list.
    for (int i = 0; i < vertexCount; i++) {
        // Print the vertex number.
        System.out.printf("%d: ", adjacencyList[i].vertexID);
        Edge edge = adjacencyList[i].edgeList;
        // Cycle through all the edges, and add to array list.
        ArrayList<Integer> edges = new ArrayList<Integer>();
        while (edge != null) {
            edges.add(edge.node.vertexID);
            edge = edge.next;
        }
        // Sort the array list.
        Collections.sort(edges);
        // Print the contents of the array list.
        for(int node = 0; node < edges.size(); node++){
            if(node+1 == edges.size()){
                System.out.print(edges.get(node));
            } else {
                System.out.printf("%d, ", edges.get(node));
            }
        }
        System.out.println();
    }
}
```

Listing 29: Solution to Adjacency List Problem (Java)

The Adjacency List Problem is a relatively simple problem that relies on the data structure implementation to be correct. As can be seen in Listing 29, within the main() function, while-loop comprises of simply reading in the vertex information for each graph, building a list in memory. The only main complication is the list being printed in an ascending manner, which is solved by creating an ArrayList of just then edges, sorting it²², and printing the result.

Sample test cases include:

- Edge count for the entire graph being zero.
- The degree of an vertex being zero.

²²Collections.sort() utilises a merge sort, so offers **O(nlogn)** performance for sorting.

A References

Throughout this guide multiple sources were used to construct this guide. The main sources of information presented through out this guide were:

- Corney, M. (2011). INB204 Special Topic Advanced Algorithms, Semester 2, 2011, Lecture Slides and Notes.
- Tang, M. (2012). INB371 Data Structures and Algorithms, Semester 1, 2012, Lecture Slides and Notes.
- Goodrich, M., Tamassia, R., (2002) *Algorithm Design: Foundations, Analysis and Internet Examples*, Wiley Press. ISBN: 0-471-38365-1 (QUT Library: 005.1 622)

Additionally, the following sections were derived from the following sources:

- Section 5.3.1 - Array Performance, Timing Values and Cache layouts, derived from:
 - Intel Corporation, (2002) *Intel Pentium 4 and Intel Xeon Processor Optimisation Reference Manual*. Intel Corporation, Order Number: 248966-007 (No longer available in print and very limited circulation, but newer editions in electronic form may be found at: <http://www.intel.com/content/www/us/en/processors/architectures-software-developer-manuals.html>).
- Section 6.1.1 - Greatest Common Divisor, Description Text, derived from:
 - Wikipedia (n.d.), Greatest Common Divisor, Retrieved August 16, 2012, from http://en.wikipedia.org/wiki/Greatest_common_divisor
 - Wolfram Mathworld (n.d.), Greatest Common Divisor, Retrieved August 13, 2012, from <http://mathworld.wolfram.com/GreatestCommonDivisor.html>
- Section 6.1.2 - Sieve of Eratosthenes, Figure 6.1, derived from:
 - Wikipedia (n.d.), Sieve of Eratosthenes, Retrieved August 15, 2012, from http://en.wikipedia.org/wiki/Sieve_of_Eratosthenes